

The Multi-Workload Challenge: Predictable Failover Across the Stack

01/22/26

Akshay Pratinav,
Staff Software Engineer, Intuit



Agenda

- Problem & Motivation
- Principles & Mental Model
- Architecture Diagram
- Outcomes & Takeaways

Problem & Motivation

Outages Are Normal



Regional failures
are inevitable



Microservices
amplify blast radius



Unpredictable recovery
causes the real damage

Why Microservices Make It Worse



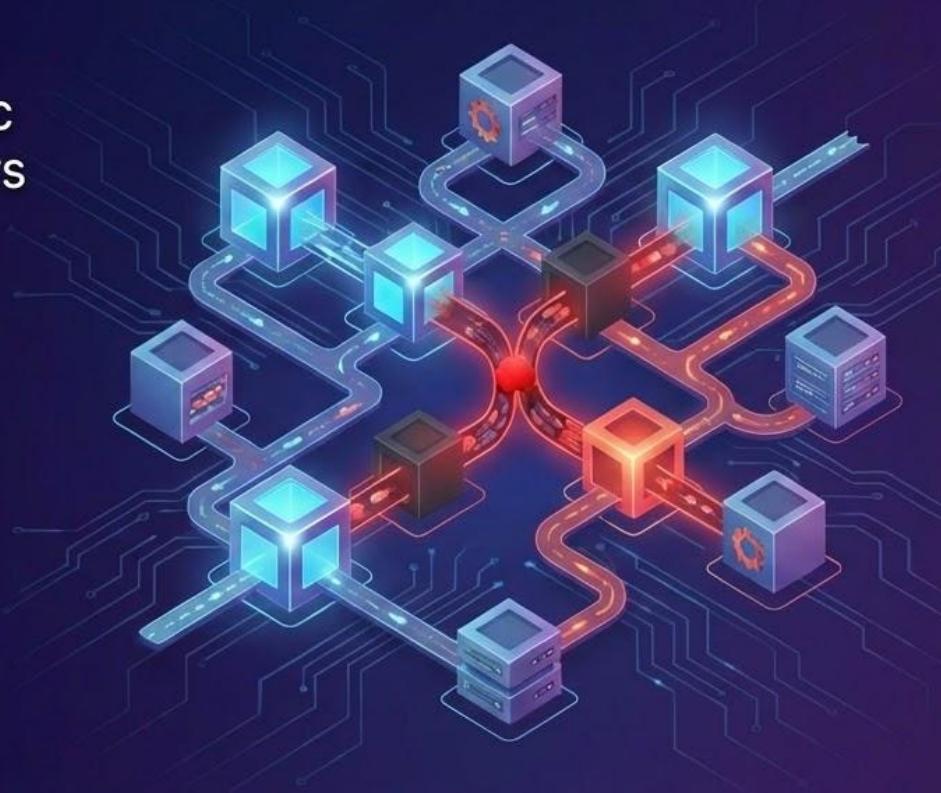
- APIs, databases, async pipelines, routing layers



- Each layer fails differently



- Recovery is gated by the slowest workload



Recovery, Not the Outage



Outages are
unavoidable

Recovery determines
customer impact

Humans under pressure
are the weakest link

The Hidden Problem



Fragmented, team-specific
failover approaches



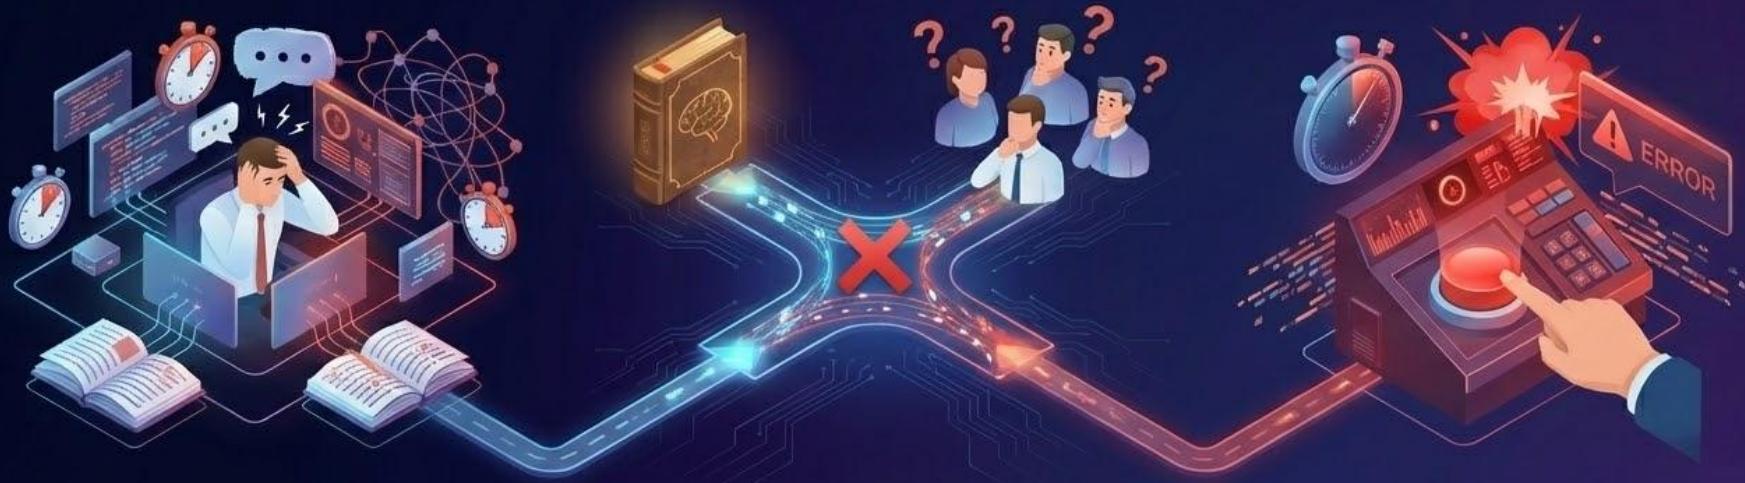
Manual scripts and
tribal knowledge



Tested only during
real incidents



Why Runbooks Don't Scale



Context switching
during incidents

Tribal knowledge risk

Error-prone manual
execution

Principles & Mental Model



What We Mean by Boring Failover



Predictable and
repeatable



Fully automated



No heroics or
guesswork



Declarative Intent

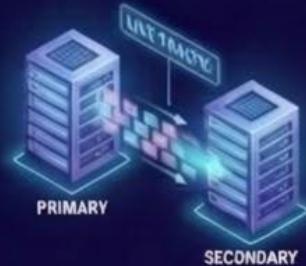
Services declare
what they need

Platform decides
how to execute

Consistent behavior
across workloads



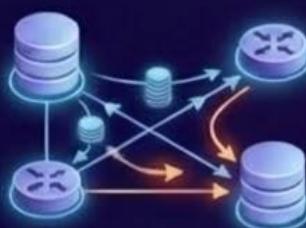
Fail Over the Entire Stack



Compute



Async and batch
workloads



Datastores and routing



Workload Diversity Matters



Different RTO / RPO requirements

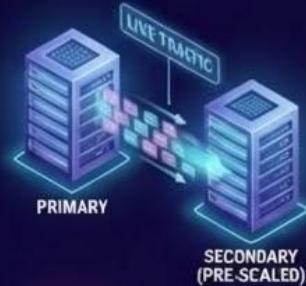


Customer-facing vs background workloads



Cost vs availability tradeoffs

Pre-Scale Capacity



Secondary region pre-scaled
using live traffic



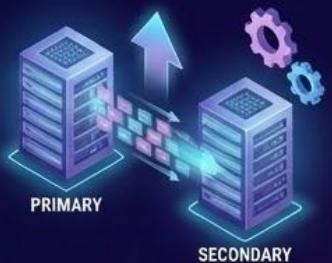
Avoid cold starts during
outages



Balance speed and cost



Data Is the Hard Part



Automated database
and cache promotion



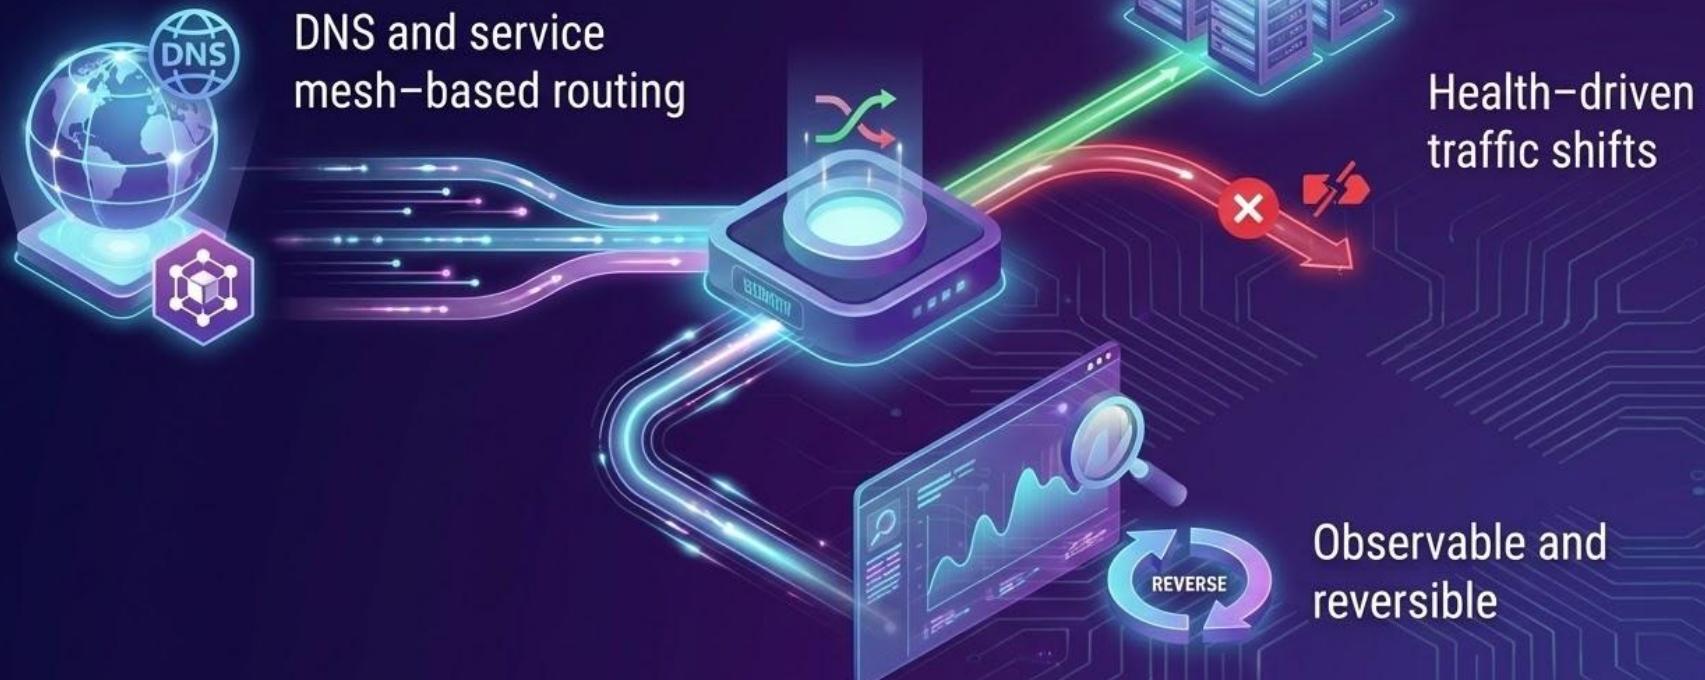
Safety and consistency
checks



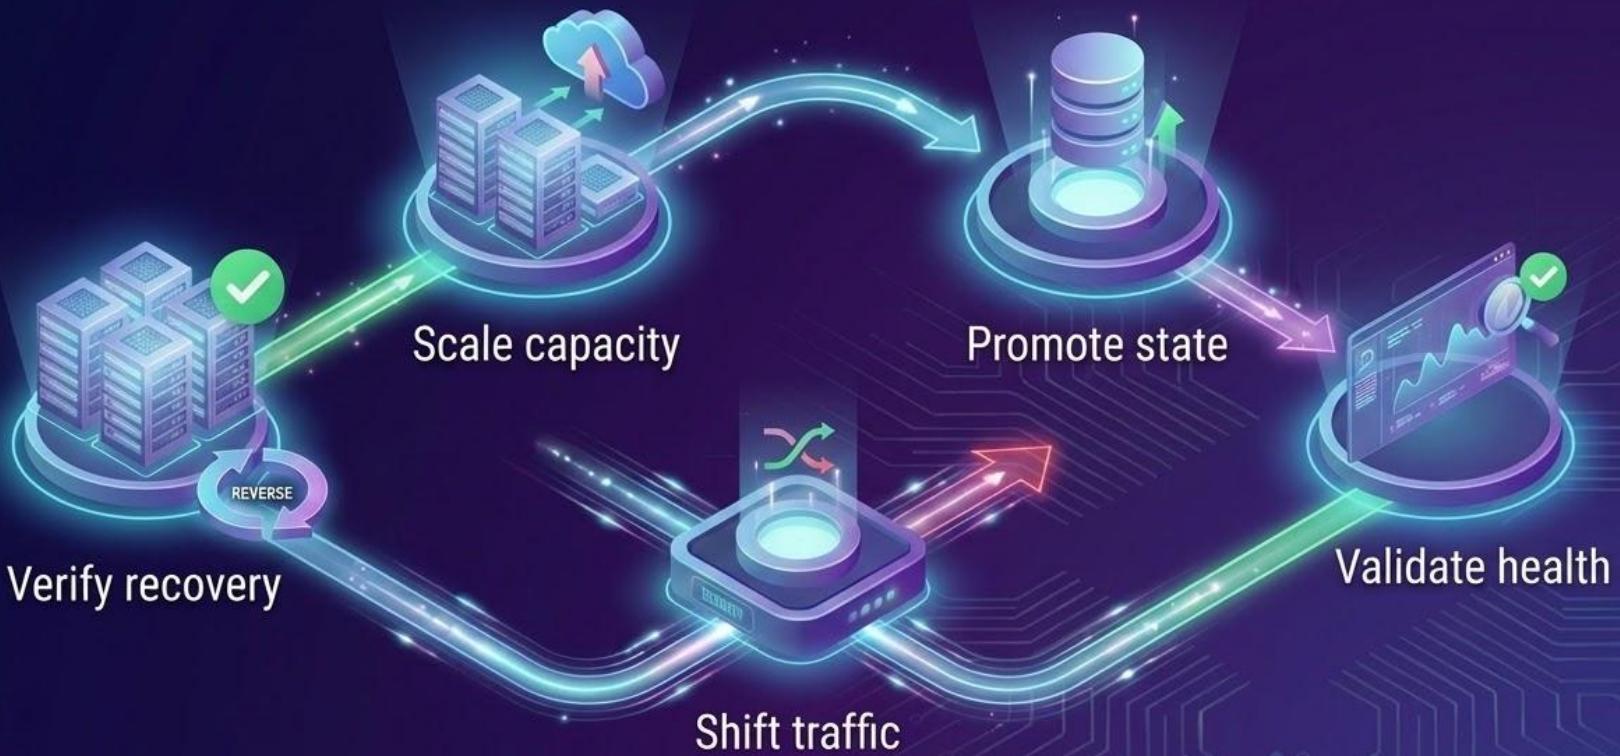
Correctness over speed



Traffic as a Control Plane



Failover Is a Workflow

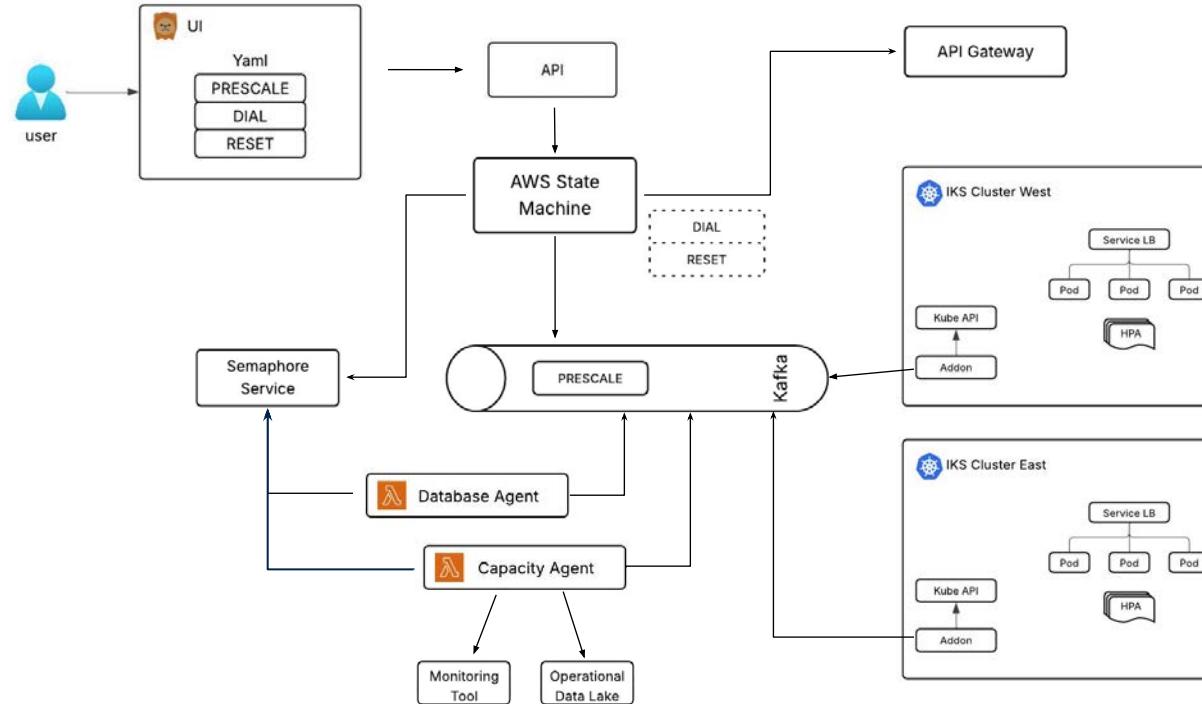


Architecture Diagram & YAML schema

DR Procedure Example

```
stages:  
  Stage 1 → prescale-east:  
    agent-id: "armador/v1"  
    capacity-compute:  
      type: 'scale'  
      cluster: 'cluster1'  
      namespace: 'namespacel'  
  Stage 2 → switch-db-primary-to-east:  
    agent-id: 'database/v1'  
    fail-over-global-clusters:  
      global-cluster-identifier: 'ewok-sample-app-global-cluster-e2e'  
      assume-role: 'arn:aws:iam::682033466980:role/database-switchover-role'  
      switch-over: true  
  Stage 3 → dial-traffic:  
    endpoints:  
    - name: 'e2e.intuit.service1.com'  
      failover-to-region: 'us-east-2'  
  Stage 4 → reset-east:  
    agent-id: 'armador/v1'  
    capacity-compute:  
      type: 'reset'  
      cluster: 'cluster1'  
      namespace: 'namespacel'
```

High Level Architecture



Safety & Operations



Guardrails Prevent Cascading Failures



Health gates before
traffic shift

Dependency
validation

Automated
rollback paths

Operational Readiness



Health gates and dependency checks



Drift detection

No special DR-only configs

Game Days Without Fear



Regular regional
failover drills



Low-stress,
repeatable execution



Confidence during
real incidents

Outcomes & Takeaways



Anti-Patterns to Avoid



Ad hoc scripts



Untested
assumptions



High cognitive load
during outages

Outcomes at Scale



Faster recovery
times



Lower operational
overhead



Predictable and
boring incidents

Cultural Shift: From Fear to Confidence



Key Takeaways



Standardize failover
across workloads



Use declarative
intent



Pre-scale
intelligently



Practice
regularly

Thank You



akshay_pratinav@intuit.com



www.linkedin.com/in/akshaypratinav