

A person in a grey hoodie stands in a server room, looking at a tablet. The room is filled with rows of server racks on both sides, illuminated by overhead lights. A digital display on the left rack shows various charts and graphs. The person is holding a tablet that also displays data visualizations.

# Observability for Enterprise AI: Reducing NLP Hallucinations with HITL

A practical framework for reliable, auditable, and production-grade generative AI systems.

**By: Arjun Danda Sureshababu**

# What We'll Cover

Why Hallucinations Are an  
Enterprise Problem

The Right Question to Ask

The Reliability Framework: A  
Layered Pipeline

Grounding, Validation &  
Human-in-the-Loop

Benchmarking & Evaluation

Six-Phase Rollout

Enterprise Checklist & Closing

# Why Hallucinations Are an Enterprise Problem



In legal review, compliance validation, privacy operations, and risk assessment, a hallucinated clause or fabricated risk score is not a curiosity it is an **operational failure**.

Large language models are optimized for fluency, not verifiable correctness. Hallucinated content reads with the same confidence as grounded content making the failure mode **dangerous precisely because it is disguised**.

# The Right Question to Ask

Not "how do we get the model to hallucinate less" but "how do we build a system where hallucination is *visible, measurable, and stoppable* before it reaches a decision-maker."

## Legal Review

Unverifiable claims carry direct liability exposure


## Compliance

Fabricated rules create false sense of coverage

## Privacy & Risk

Outputs must be defensible after the fact, not just plausible

# The Reliability Framework: A Layered Pipeline



Each layer produces its own checkpoint and signal turning a black-box model call into an auditable assembly line where failures can be precisely located and corrected.

# Pre-Generation Guardrails

## Prompt Templates as Contracts

Versioned, tested templates specifying task boundaries, permitted sources, output format, and behavior when information is missing.

## Instruction Constraints

Explicit behavioral rules: distinguish "stated in source" from "inferred," never present inference with the same confidence as a sourced fact.

## Role Enforcement

Narrows the model's persona to the task preventing "scope creep" hallucinations where the model wanders into unsupported territory.

# Grounding Through Hybrid RAG

## Why Hybrid Retrieval?

### Sparse (BM25)

Exact term matching — fast, interpretable, ideal for defined terms and clause numbers. Best when query vocabulary matches the source exactly.

### Dense Vector Search

Semantic similarity retrieves relevant passages even when vocabulary differs — essential for heterogeneous enterprise document sets.

### Reranking & Multi-Query

Reranking elevates the most relevant passages before generation. Multi-query orchestration decomposes complex questions into sub-queries, covering multiple retrieval targets simultaneously.

# The Observability Payoff of Retrieval

Every generated output can be logged alongside the exact passages retrieved to support it turning "was this hallucinated?" from a subjective judgment into a **checkable claim**.



## Paper Trail

Every output has a traceable source no comparison baseline is lost



## Auditable Claims

Reviewers can directly compare generated text against retrieved source material



## Defensible Output

Grounded outputs hold up to regulatory and legal scrutiny after the fact

# Post-Generation Validation

1

## Citation Enforcement

Every factual claim must reference a specific retrieved passage. Uncited claims are routed for mandatory review.

2

## Schema Validation

Outputs for structured use cases (risk scoring, checklists) are checked programmatically malformed responses are flagged instantly.

3

## Self-Critique

A second model pass evaluates the output against retrieved context, filtering low-confidence claims before human review.

4

## Multi-Step Verification

Higher-stakes outputs chain multiple verification passes each producing its own log entry, turning failures into improvement opportunities.

# Human-in-the-Loop: Structured Expert Review

Rather than ad hoc thumbs-up/thumbs-down, HITL uses **structured rubrics** scoring outputs across four independent dimensions:

## Factual Correctness

Alignment with source material and known facts

## Contextual Relevance

Does the output address the specific question — not a generic adjacent one?

## Hallucination Severity

How consequential is unsupported content when present?

## Tone Alignment

Appropriate confidence level — avoiding false certainty or unhelpful vagueness

# Sampling, Escalation & Feedback Loops



## The Feedback Loop

Reviewer scores flow directly into prompt and guardrail refinement, while annotations recalibrate automated validator thresholds to reduce future errors. Each review decision expands a labeled dataset that powers ongoing evaluation and fine-tuning — making the entire system **measurably better over time.**

# Benchmarking Under Stress

## Adversarial Prompts

Tests whether the system degrades gracefully or produces fluent, confident, wrong answers under injection attacks.

## Missing-Knowledge

The key diagnostic: does the system admit ignorance or fabricate a plausible-sounding answer?



# HUMBR: A Better Evaluation Approach

## The Problem with Single-Reference Evaluation

Comparing outputs against one human-written "ground truth" penalizes equally valid, differently phrased, or more complete responses.

## The HUMBR Solution

**Hybrid Utility Minimum Bayes Risk** evaluates candidates against a *distribution* of plausible good answers rewarding outputs that are robustly useful across the space of acceptable answers, not just similar to one reference.

## Evaluator Preference: Clear Majority

In comparative testing, HUMBR-selected outputs were **preferred by human evaluators** over human-crafted ground-truth answers in a clear majority of head-to-head comparisons.

This makes automated quality monitoring **less brittle** and better aligned with what actually matters: defensibility across multiple qualified reviewers.

# Six-Phase Rollout

## Phase 1

Establish guardrails & scope

## Phase 2

Build grounding layer

## Phase 3

Add post-generation validation

## Phase 4

Stand up HITL review

## Phase 5

Build benchmarking program

## Phase 6

Instrument for observability

# Enterprise Rollout Checklist

## Guardrails

Prompt templates versioned as controlled config with explicit, falsifiable role boundaries.

## Validation

Citation enforcement, schema validation, self-critique, and multi-step verification in place.

## Feedback Loops

Reviewer input feeds back into prompt refinement and validation thresholds on a defined cadence.

## Governance Ownership

Named owner accountable for reliability metrics with a defined degradation response process.

# Building AI Systems That Are Genuinely Trustworthy

Hallucination mitigation is not a problem solved once at the model layer. It is an **ongoing systems and governance challenge**.

## Observable

Every layer inspectable, logged, and auditable independently

## Measurable

Reliability tracked over time not just reacting to individual incidents

## Improvable

Each layer independently tunable as models, prompts, and data evolve

A system that cannot explain why it produced a given answer is **not ready for high-stakes production use** regardless of how fluent its outputs appear.

# Thank You

For questions or follow-up, feel free to reach out.

**Arjun Danda Sureshababu**