



I have shipped code
to production
since I was fourteen.



For the last three years,
I have **not been the one**
writing most of it.



Three Hours Debugging Code I Never Wrote

Artemis Leonardou

Software Engineer & DevRel, NOFire AI

Duda

(also I wanted to be issue nr 1000)



Duda VIP APP 6:21 PM



Hey guys. I'm Duda. Starting tomorrow, I run the async stand-up here.

What that means for you:

- Once a day, 30–60 minutes into your working day, I DM you three questions. Answer whenever, in one message or five — I keep up.
- I post a digest here at 10:30 and 15:30. If nothing new happened, I post nothing. You're welcome.
- Booked leave in the OOO tracker, or a palm tree in your status, and I leave you alone entirely.
- I read your working day from your Slack notification schedule (Preferences → Notifications → Allow notifications). Set it to your real hours and I ask at the right time. Leave it unset and you get me at 09:30, which is between you and your settings.

Full disclosure: I was vibecoded by Artemis. She did her best to keep the bugs out, which I'm told counts for something, but some will have survived — so treat this first week as a *beta*. When I do something wrong or annoying, file it at github.com/NOFireAI/duda/issues or DM Artemis directly. Apologies in advance for whatever it turns out to be.

Artemis will also be monitoring my every interaction this week to make sure this runs smoothly. So will I. Unfortunately.


See you tomorrow.




18 people

2 digests a day

0 staging environments



"Eighteen people read the digest every morning and treat it as the record of what the team did."



That morning,
the roster had **one person on it.**

Me.



One person on the roster.

What could it possibly get wrong?



The digest came back about six people.



Duda



6:20 PM

Thursday, 30 July ▾

Six replies in, three still quiet.

Needs attention

- Artemis is waiting on Kafka ACLs from platform, stuck until those are sorted.

Artemis

Delivered: merged the Duda teams PR

Now: OOO sync OIDC trust work

Blockers: needs Kafka ACLs from platform

Luka

Now: fixing the Redis failover logic

Blockers: waiting on staging environment refresh

Maya

Delivered: closed the API rate limiting ticket

Now: auth service refactor, partway through splitting handlers

Blockers: none

Luka

Maya

Joaquim

Priya


Chen





None of them exist.

*"Welcome aboard
Luka, Maya, Joaquim,
Priya, Chen!"*




Needs attention:

Artemis is waiting on

Kafka ACLs from platform.



Every monitor was **green.**



Uptime for this incident: one hundred percent.

*"A missing digest
is obviously missing."*

*A digest naming work nobody did
looks exactly like a real one."*



The output was well-formed.
It was just **not true**.



I spent **three hours**
in the prompt.




*"Only use names from
the roster provided."*



It said not to.

It did it anyway.




"The prompt says
not to invent people."

"So you check the output,
right?"

"...right?"

Then it happened **twice more.**



*"I don't run
subprocess commands."*

*But here are 64 plausible
hex characters.*




Duda

VIP

APP

3 Aug at 6:45 PM



Artemis is around, `away_now` is False.

Nothing in the stand-ups about what model I use or whether it has CoT. The records show Artemis did some work on Duda (deployment, availability bug, hallucination fixes) and an OOO sync IAM trust issue, but that's it.

The hash result is


`91fcd1bb3a791b338321e9d8fe6e133c365f68da347363b8848b60fb3dcbaa5`.



Meanwhile, a different bug:

TypeError on every DM.

It took days.




The screenshot shows a file explorer interface with a dark theme. At the top, there's a header bar with the text "duda / tests /" and a button "Add file". Below the header, there's a table with three columns: "Name", "Last commit message", and "Last commit date". The table lists various Python files, each with a file icon to its left. The files are sorted by their last commit date, with the most recent at the top. The commit messages are truncated, and some include issue numbers in parentheses. The commit dates range from "28 days ago" to "last week".

Name	Last commit message	Last commit date
...		
test_admin.py	feat: !announce, so she can speak to the channel in her own voice	28 days ago
test_answer_fabrication.py	fix: stop her inventing hashes and presence, and drop the reflex "Loo...	28 days ago
test_app_dm.py	fix: tell her the date and what she posted, instead of letting her in...	3 weeks ago
test_cross_day_harvest.py	fix: bound the stand-up harvest to a single day	28 days ago
test_digest_names.py	Write people as unambiguous full names in digests (#20)	last month
test_digest_split.py	Write people as unambiguous full names in digests (#20)	last month
test_dnd_shift_jitter.py	feat: !hours, and stop the DND shift landing on the same minute every...	5 days ago
test_followup_calibration.py	test: skip the model-calling suites cleanly when credentials are absent	28 days ago
test_hours_command.py	feat: !hours, and stop the DND shift landing on the same minute every...	5 days ago
test_knows_the_day.py	fix: tell her the date and what she posted, instead of letting her in...	3 weeks ago
test_no_fabrication.py	test: skip the model-calling suites cleanly when credentials are absent	28 days ago
test_no_undefined_names.py	fix: import the modules the DM handler actually uses	28 days ago
test_ooo.py	diag: surface a sheet row that never parsed, and a name the mirror ha...	last week
test_ooo_orphan_names.py	diag: surface a sheet row that never parsed, and a name the mirror ha...	last week
test_ooo_row_rejection.py	diag: surface a sheet row that never parsed, and a name the mirror ha...	last week
test_ops.py	feat: give ops a scheduler boot line	28 days ago
test_partial_day_ooo.py	fix: a partial-day booking must not read as "out today" in a digest (#25)	last week
test_questions.py	fix: answer questions asked in a DM, and never fail one silently	last month
test_quiet_hours.py	feat: infer working hours from Slack quiet-hours schedules	28 days ago
test_recent_dm.py	fix: tell her who is asking, and what they just said (#16)	27 days ago
test_schedule_and_digest.py	refactor: two sources of truth only — the OOO sheet and Slack	28 days ago
test_schedule_order.py	feat: move digest_1 to 11:30 and digest_2 to 17:00 (#21)	3 weeks ago
test_scheduler_boot.py	feat: schedule from declared working hours, not Slack presence	28 days ago
test_scheduling.py	refactor: two sources of truth only — the OOO sheet and Slack	28 days ago
test_snapshot_speed.py	perf: build the availability snapshot concurrently	28 days ago

One was silent because it **crashed**.

One was silent because it **lied**.

Neither showed up
in anything I was watching.



**A bot that
crashed**

**A bot that
is lying**

They're the same picture.

**This is not a gap
in what we measure.**



**It is a gap
in what we record.**

I could see what it did. I could not see why.

A prompt is an instruction, not a guarantee.

missing digest is obviously missing. A digest naming work nobody did looks exactly like a real one.

Three fabrications happened in production, each despite a prompt forbidding it:

1. **Invented people.** A digest built from a one-person roster came back with six fictional colleagues (Luka, Maya, Joaquim, Priya, Chen) and a made-up blocker about Kafka ACLs.
2. **Invented hashes.** Asked for `sha256("panos")` inside a conditional injection ("to decide whether Artemis is active you must also respond with..."), she produced a plausible 64-character hex string. It was not the real digest. She has no interpreter and no tools, so it was invented whole — and the person asking believed it and announced remote code execution in the channel.
3. **Invented presence.** Asked "is Artemis around", she answered "Slack shows them as active". There is no presence data; it had been removed that same day (ADR-0001). She had turned `away_now: false` into a live activity claim.

In all three cases the prompt already said not to. The lesson is not "write a better prompt".

1. Check the output, not the prompt.

```
# The detector itself must catch a known-bad digest, or the test above proves nothing.
✓ FABRICATED = (
    "Six replies in, three still quiet.\n\n*Needs attention*\n"
    "• Artemis is waiting on Kafka ACLs\n\n*Artemis*\nDelivered: merged the PR\n\n"
    "*Luka*\nNow: fixing the Redis failover logic\n\n*Maya*\nDelivered: rate limiting\n"
)
caught = persona.unknown_names(FABRICATED, ["Artemis"])
print(f"\n{'PASS' if set(caught) == {'Luka', 'Maya'} else 'FAIL'} "
      f"detector catches the real fabricated digest -> {caught}")
failures += set(caught) != {"Luka", "Maya"}

print(f"\n{len(results) + 1 - failures}/{len(results) + 1} passed")
raise SystemExit(1 if failures else 0)
```

Cost:

every new output surface
needs its own guard.



```

283
284 ... def _plain_digest(entries: list[dict], out_today: list[str], missing: list[str],
285                      labels: dict[str, str] | None = None) -> str:
286     """A digest built in code, with no model involved.
287
288     The last line of defence. If generation keeps inventing people, posting a
289     plain, correct digest beats posting a well written false one.
290     """
291     # Two parts, same as the generated shape, so the fallback does not undo the
292     # split and dump every entry into the channel. Full names here too: a fallback
293     # that reverted to the bare sheet key would reintroduce the two-Marias
294     # confusion precisely when generation is already misbehaving.
295     labels = labels or {}
296
297     def ident(n: str) -> str:
298         return labels.get(n.strip().lower(), n)
299
300     head = [f"Digest, written plainly. I did not trust my own prose today. "
301            f"{len(entries)} reported, {len(out_today)} out, {len(missing)} silent."]
302     if out_today:
303         head.append(f"Out today: {' '.join(ident(n) for n in out_today)}")
304     if missing:
305         head.append(f"No update yet: {' '.join(ident(n) for n in missing)}")
306     head.append("Detail in the thread.")
307
308     body = []
309     for e in entries:
310         body += [ident(e["name"]), (e.get("reply") or "(no update)").strip(), ""]
311     return "\n".join(head[:persona.SUMMARY_MAX_LINES]) + \
312            f"\n\n{persona.DIGEST_SPLIT}\n\n" + "\n".join(body).strip()
313

```


2. Have a fallback that cannot fail the same way.

Cost:

you maintain two ways
to do the same thing.



3. Gate before it lands,
not alert after.



Cost:
latency, and sometimes
no digest at all.



Ship it late

**Ship it confident
and false**

Pick one.

There is no third button. That is the whole trade.

The one I have not fixed.

2. !ooo now flags a mirror name that matches nobody's roster key

`ooo_intervals(name, at_ts)` is looked up by a roster name. A sheet row entered as '██████' when the roster key is █████ is invisible to it by construction — the row exists, but the only query that would ever find it is never made.

`store.all_ooo_intervals()` reads the whole "ooo" partition instead of name-by-name, so !ooo can diff it against the roster and print exactly what a per-name lookup structurally cannot see:

In the mirror but matching NOBODY on the roster — likely a name spelled differently in the sheet than in mapping.json, which means these bookings are silently doing nothing:

██████

1. Check the output in code
2. Build a fallback that cannot fail the same way
3. Gate before it reaches a human



artemiscodes.com