


Enhancing Operational Efficiency with Observability

Discover how observability transforms complex distributed systems by providing deep insights into application behavior, system health, and user experience. This presentation reveals practical strategies for implementing observability to dramatically improve reliability, accelerate troubleshooting, and optimize performance across your entire technology stack.

By: Hardik Patel





Hardik Patel

Director of Engineering - Data Platforms

- 20+ years in tech industry (small startups to large enterprises)
- Masters in Software Engineering from BITS, Pilani
- Held various roles from SWE to TL to Architect to Manager to Director
- Technical Leadership at high-tech places Yahoo!, Symantec, Groupon, PayPal
- Led varied sizes of teams from 5 to 30+ engineers, leads & managers
- 10+ years of designing, developing Distributed Systems

The Challenges of Distributed Systems

Increased Complexity


Modern architectures contain hundreds of interconnected services. Traditional monitoring falls short.

Rapid Change

Continuous deployment creates constantly evolving systems. Pinpointing issues becomes harder.

Scale Issues

Problems may only appear at scale. Test environments can't replicate production production loads.



The Three Pillars of Observability



Metrics

Numerical data points collected at regular intervals, revealing system performance patterns and health. Enables identification of trends, anomalies, and potential bottlenecks before they become critical.



Logs

Timestamped records of discrete events and actions within a system. Provides rich contextual information for debugging, forensic analysis, and understanding the sequence of incidents.



Traces

End-to-end journey tracking of requests as they flow through distributed services. Illuminates complex interactions, latency issues, and dependencies that impact overall system performance.

Metrics That Matter



System Latency

Measure response times across service boundaries. Track p95 and p99 percentiles.



Error Rates

Monitor failed requests and exceptions. Set thresholds based on business impact.



Resource Utilization

Track CPU, memory, and network usage. Prevent resource exhaustion before it occurs.



User Experience

Measure real user metrics. Connect backend performance to frontend experiences.





Distributed Tracing in Action



Frontend Request

User interaction initiates a traced HTTP request with unique correlation ID

Gateway Service

Propagates trace context while routing to downstream microservices

Database Query

Span indicates performance bottleneck with 200ms latency
latency exceeding threshold

Response Chain

Completed trace shows full request path with timing metrics
metrics across services



Business Impact of Observability

60%

Faster Troubleshooting

Reduction in time to identify root causes

25%

Improved MTTR

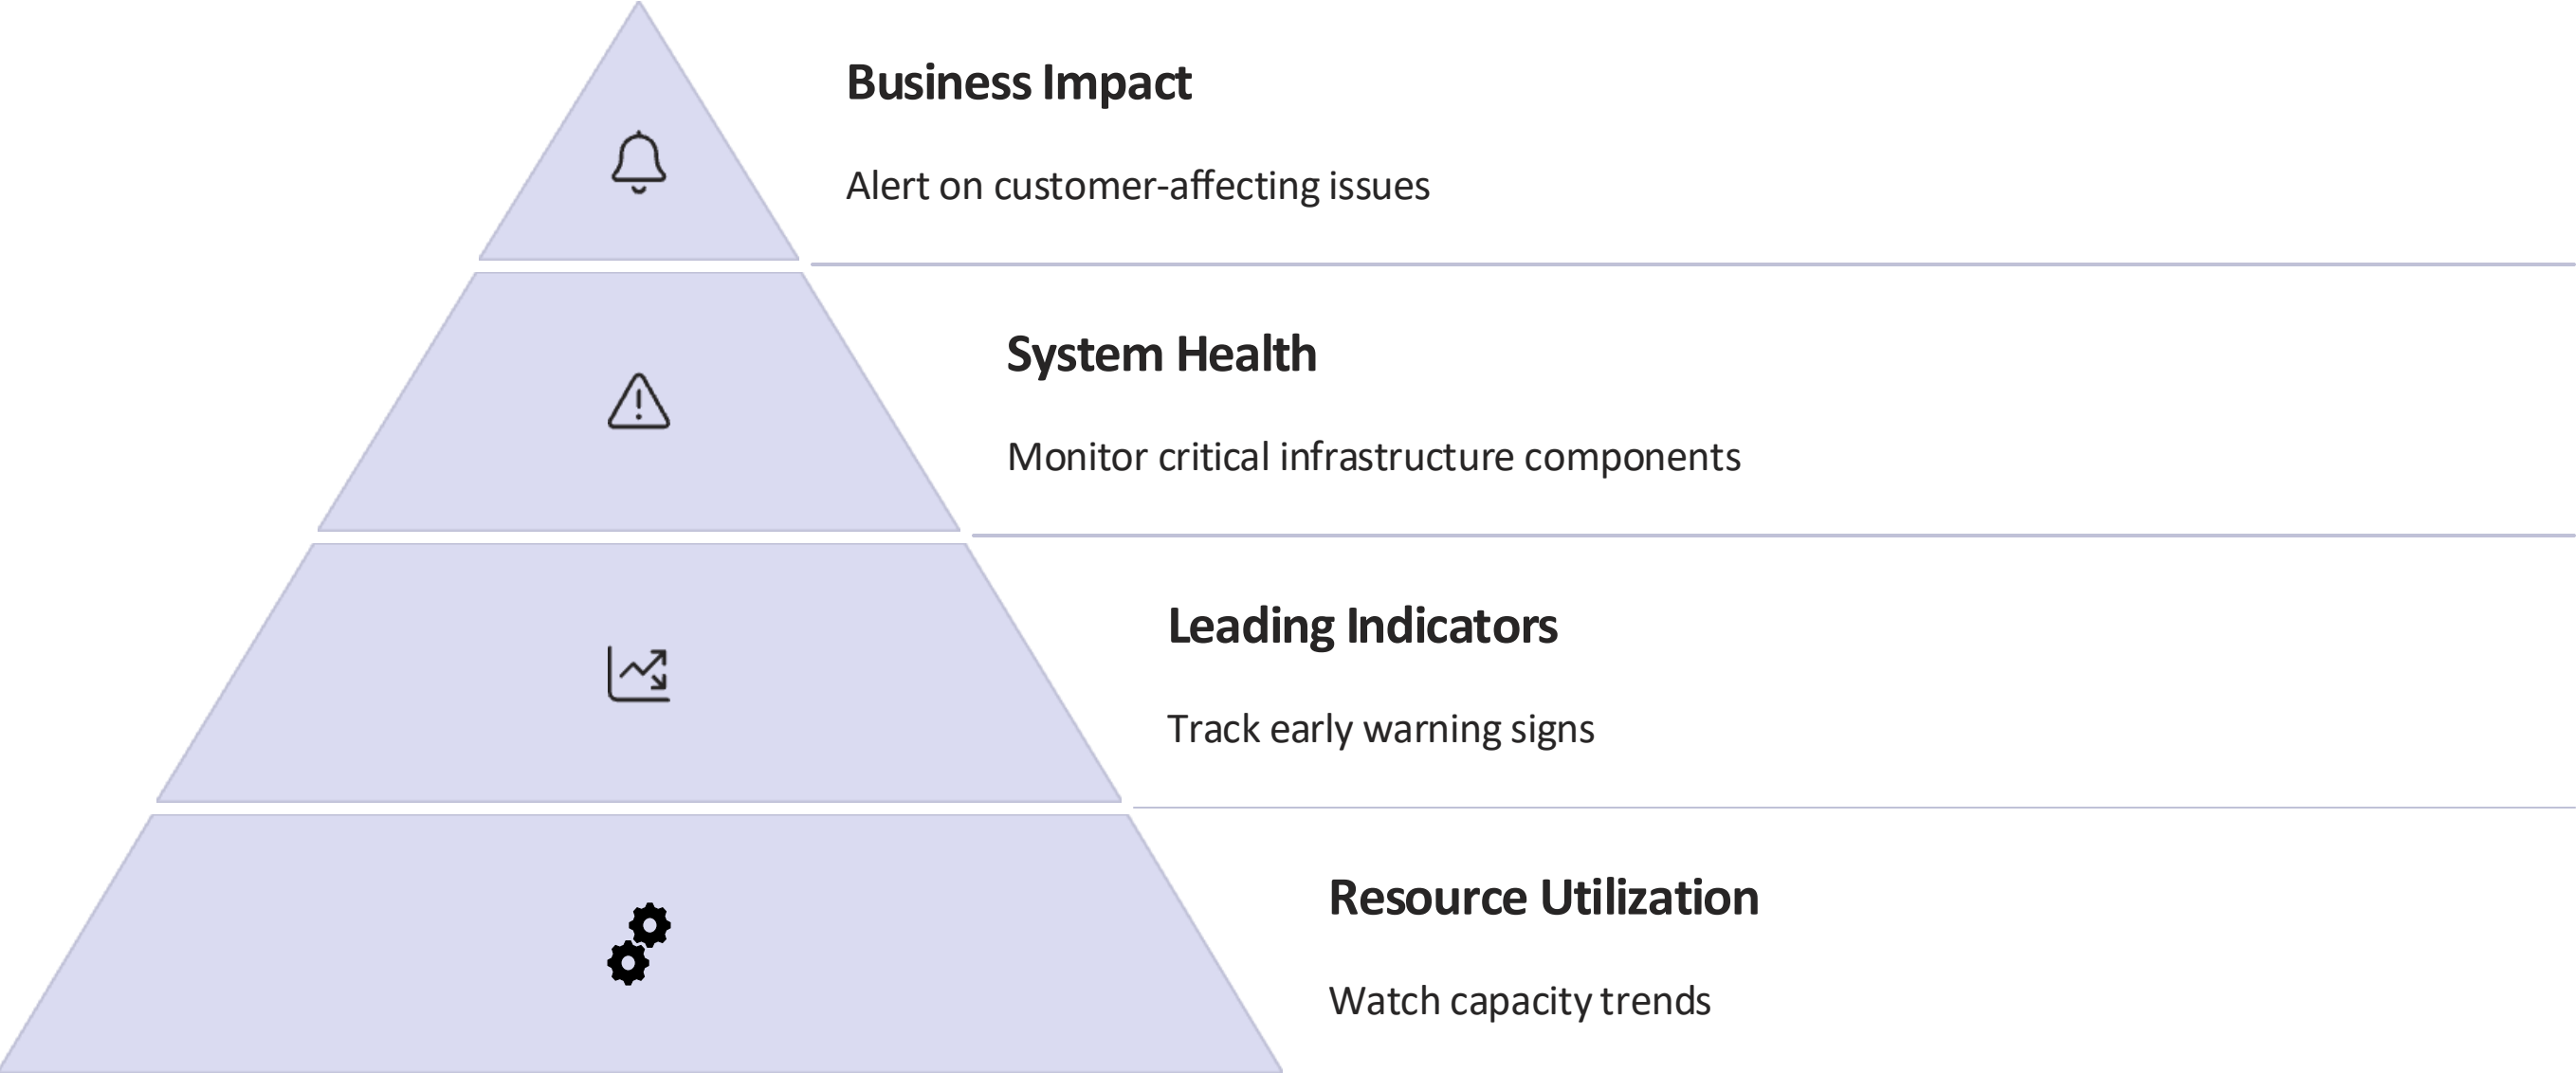
Faster mean time to recovery after incidents

30%

Increased Uptime

More reliable systems with fewer outages
outages

Implementing Effective Alerting



Observability Implementation Roadmap

Define Key Metrics

Identify the critical signals that indicate system health. Focus on business outcomes. outcomes.

Instrument Code

Add metrics, logging, and tracing to applications. Use consistent naming conventions.

Centralize Data

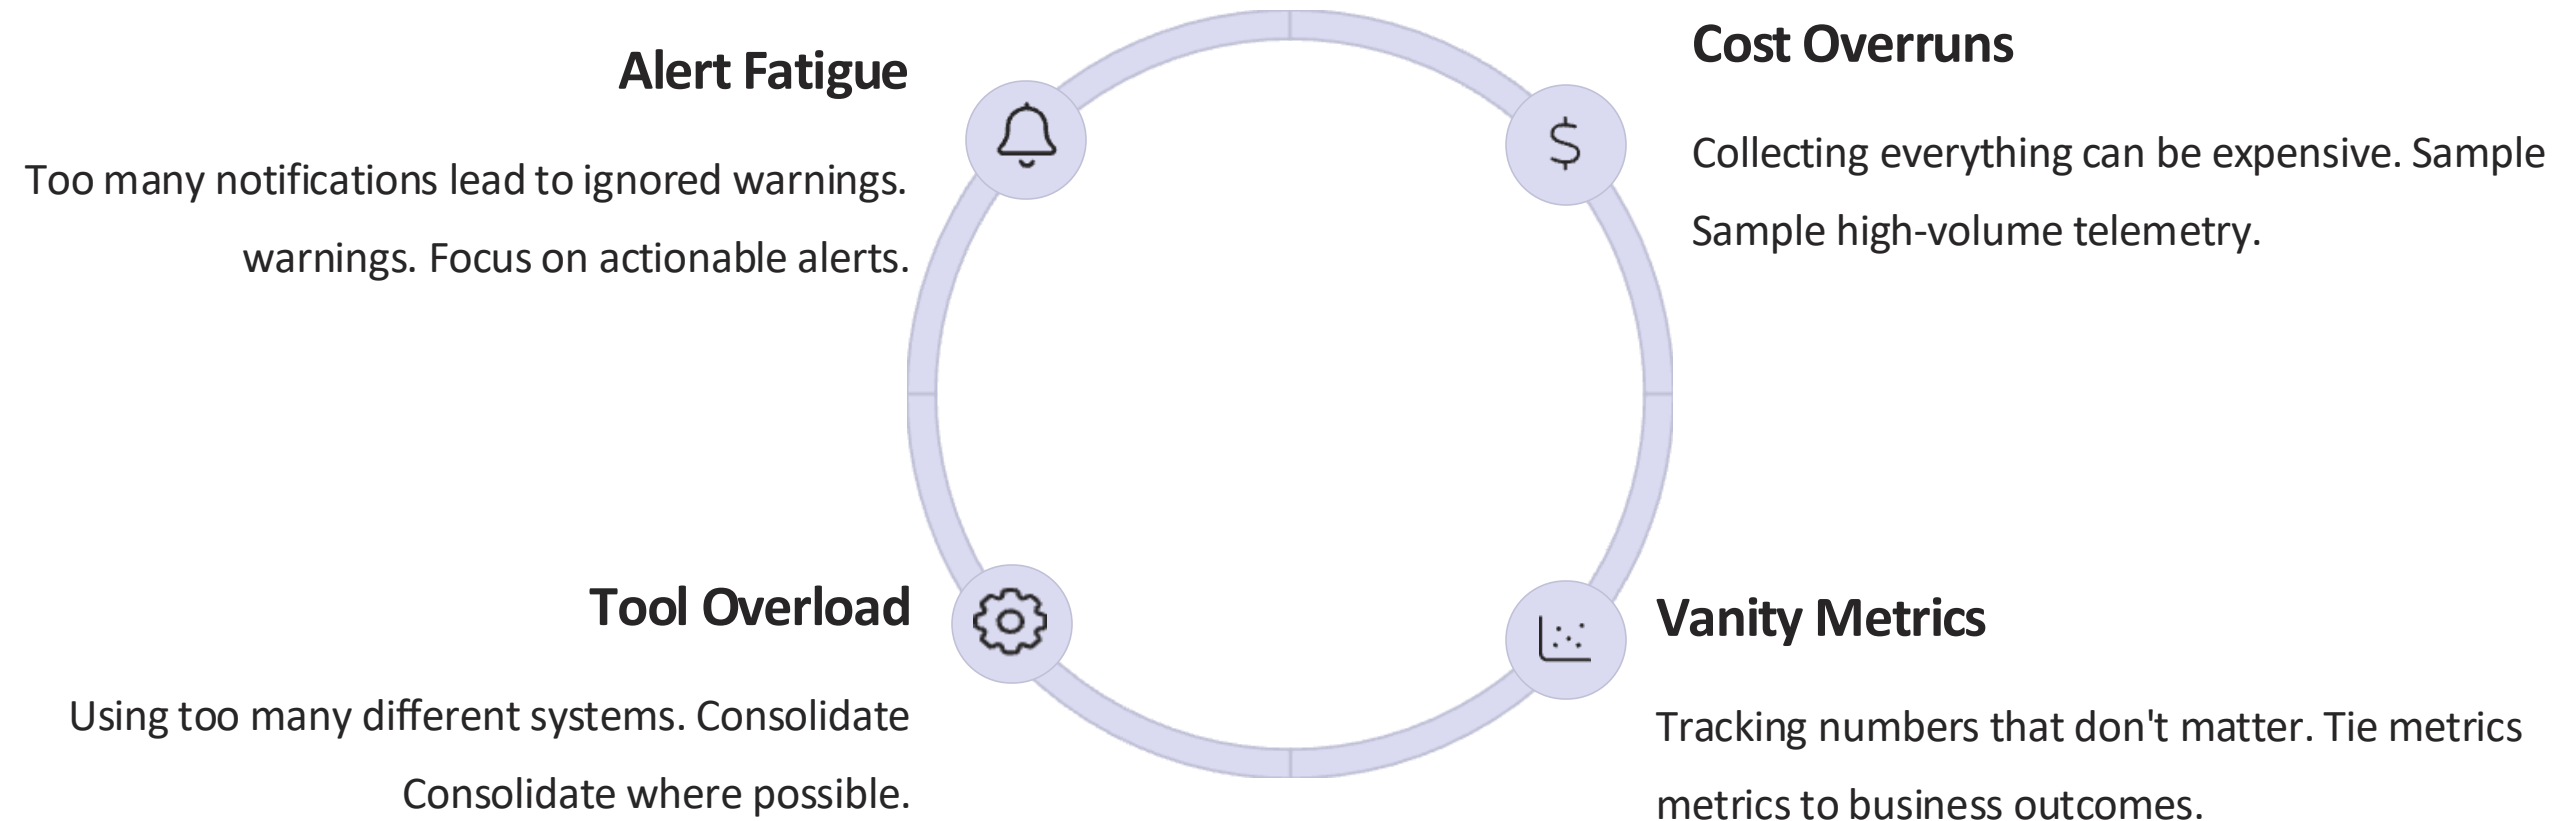
Implement collection and storage solutions. Enable correlation across services. services.

Visualize & Alert

Create dashboards and alerting rules. Eliminate noise and false positives. positives.



Common Pitfalls to Avoid



Key Takeaways



Observability is essential

Not optional for modern distributed systems



Integrate all three pillars

Metrics, logs, and traces work together



Measure business impact

Connect technical metrics to business outcomes



Build observability culture

Train teams to leverage these insights

Thank You!