


# **Designing a Read-Intensive Architecture for Information Systems with High Data Throughput**

**Building high-load systems isn't just adding resources. Efficient storage, caching, and network optimization are key, while well-designed data schemes and architectural decisions ensure performance and SLA compliance at scale**


**Ivan Sinitin**

# What does a typical Read-Intensive system look like

- K8S
- Microservices
- Sharded cache with replication
- Sharded NoSQL DB with replication




# When "Just Add Resources" Breaks



- Tail Latency (p99)
- Entropy of protocols (Gossip)
- Contention (The struggle for resources)
- Throughput vs Latency
- High Cardinality and Drift Topology

# Tail Latency


- Fan-out scheme: one API request throttles 50 shards
- Out of 1,000 nodes, the probability of encountering a slow container is 100%. One slow shard slows down the entire response.



# Entropy of protocols


Scalability Tax: Increased overhead with increasing number of nodes

- Orchestration
- CPU
- Network traffic




# Contention & Resource Exhaustion

- N service pods
- M cache pods
- $M * N$  connections
- $N * M$  descriptors
- ...




# Throughput vs Latency



# Defense Line: "Painkillers" for Highload

- Graceful Degradation: A cache crash doesn't equal a service crash (enable Circuit Breaker)
- Soft Timeouts: Don't wait forever for database. 200ms—and we'll serve what we have
- Stale-data: "Expired" is better than a 500 error

# Zonal isolation (Multi-AZ) and its price



# The Illusion of Stability: Why Symptoms Are Hidden, But the Disease Remains

- Hidden Tails: Your Soft Timeouts simply "truncate" the graph, turning real database bottlenecks into "old data" for the user
- Entropy of protocols: In Multi-AZ, Database service traffic between zones is growing exponentially
- Resource Death: 1,000 pods still hold 50,000 connections to Cache. Circuit breakers do not reduce the socket count

# Optimization Strategy: Where are we losing resources?



## Process data

Inefficient data format and schema



## Storing data

Inefficient data format and schema



## Transfer data

Inefficient data formats  
Poor understanding of data format encoding schemes

# Storage: More than just bytes

- JSON: Simple, but it's a text format. Expensive to parse, redundant
- Protobuf: Binary standard. Compact, strict, efficient
- FlatBuffers: Zero-copy access. Read only what's needed, without deserializing the entire object
- Columnar approach (Avro/Parquet): Read only the required fields from wide data structures
- Schema Registry: Extract metadata from the message body


# When convenience gets in the way of performance

	AoS (Array of Structures)	SoA (Structure of Arrays)
schema	<pre>message Request {     message Item {         int64 id = 1;         float price = 2;     }     repeated Item items = 1; }</pre>	<pre>message Request2 {     repeated int64 ids = 1;     repeated float prices = 2; }</pre>
use	easy	hard
aggregation	poor memory locality	good memory locality
runtime effects	<ul style="list-style-type: none"><li>- large number of objects in heap</li><li>- slow serialization</li></ul>	<ul style="list-style-type: none"><li>- low number of object in heap</li><li>- fast serialization</li></ul>
SIMD	no	yes

# Architectural Optimization: “CQRS” and Prepared Data

Use process-efficient


- Data format
- Data schema
- Data storage



# Caching at your fingertips




Add K8S-pod caching



# Pod sharding


## Use

- Service discovery
- Client-side balancing
- Consistent / Rendezvous hashing





# Choose different formats for different storages

- For external storage choose
  - backward compatibility
  - size-efficient
- For pod (in-memory storage) choose
  - high-performance / low resources
  - size-efficient



# Economy and Risk: Density vs. Stability

- Advantages
  - less traffic
  - more resource-efficient
  - less pods / nodes you need
- Disadvantages
  - complex system



# Life with compromises

- Team Cognitive Load
- Operational Complexity
- Invalidation and Consistency
- Risk of Cascading Collapse

# **Designing a Read-Intensive Architecture for Information Systems with High Data Throughput**

**Building high-load systems isn't just adding resources. Efficient storage, caching, and network optimization are key, while well-designed data schemes and architectural decisions ensure performance and SLA compliance at scale**

**Ivan Sinitin**