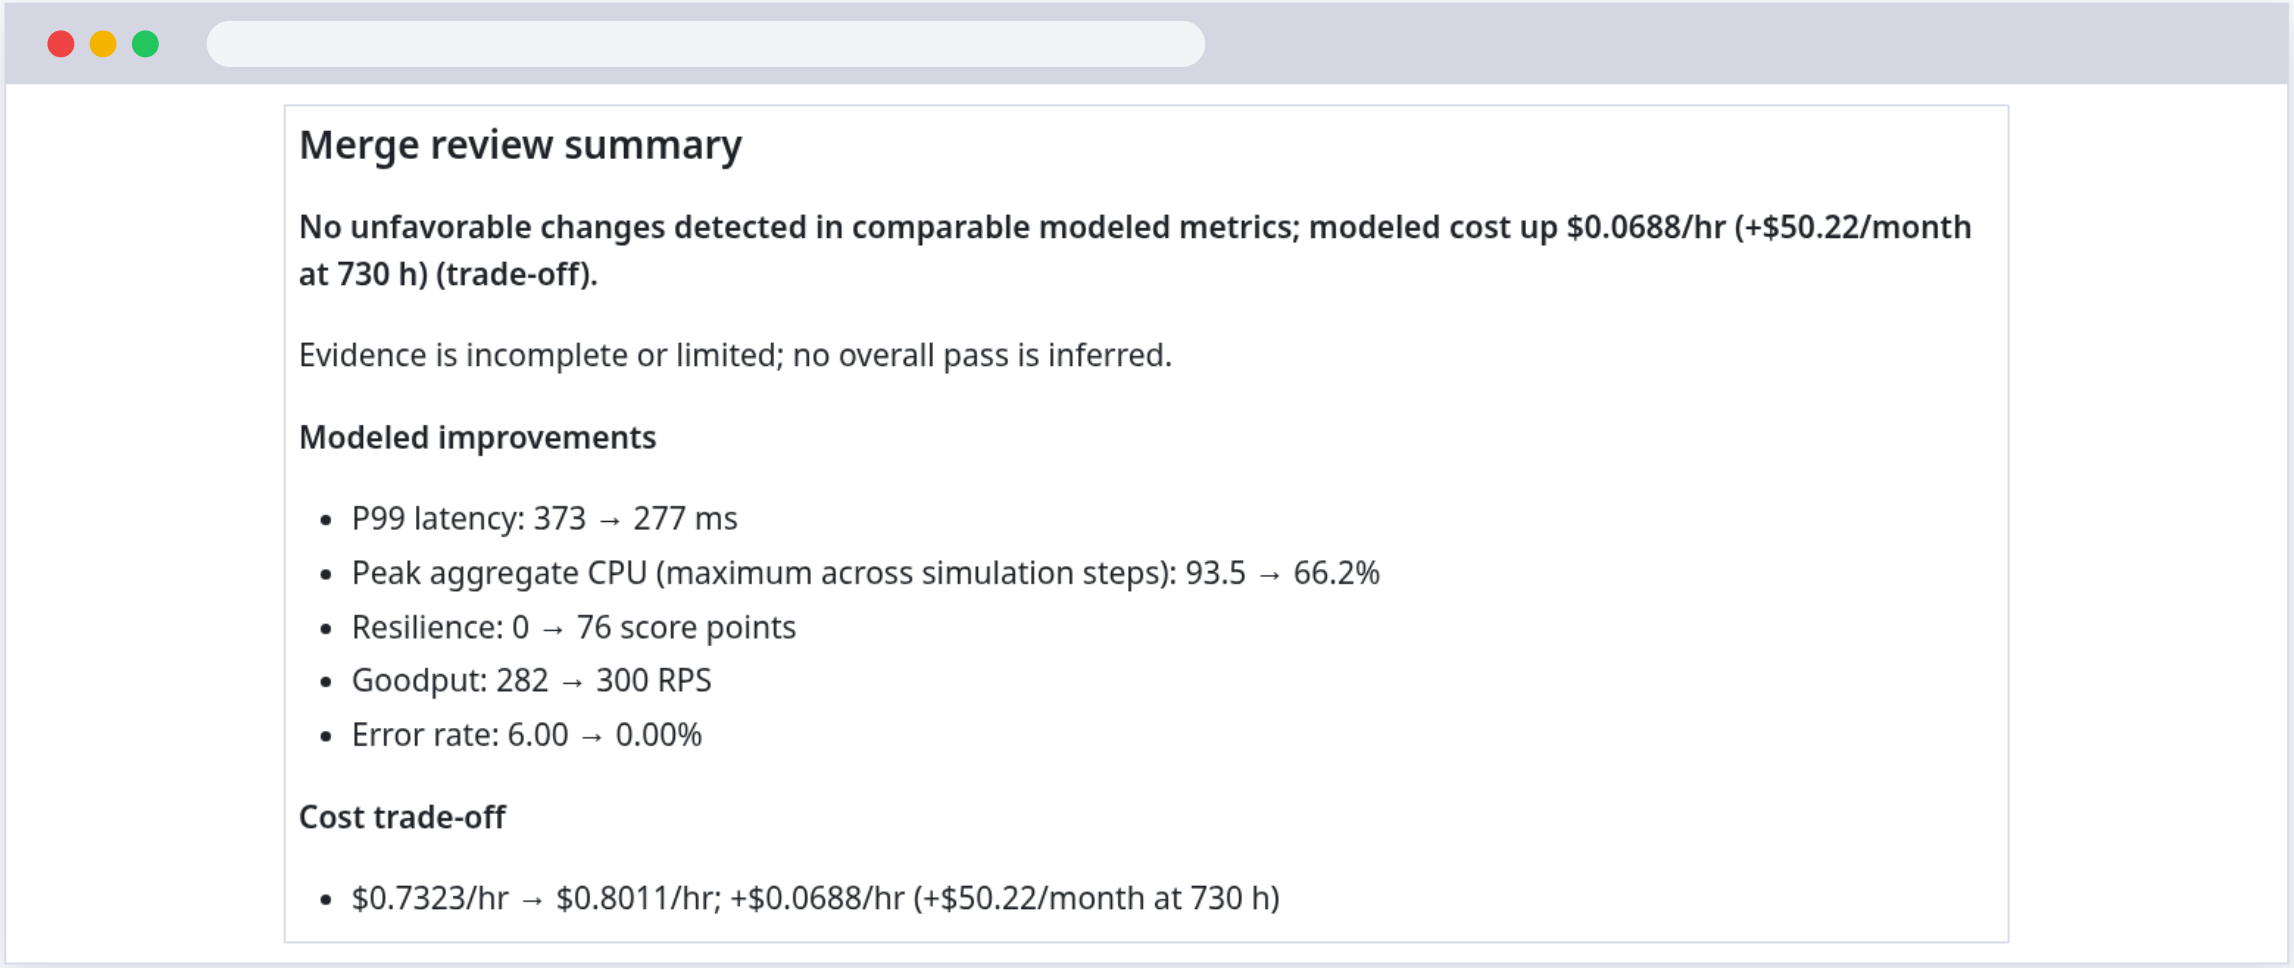
Your Agent Can Write Terraform.
It Can't Feel What It Breaks.



What the agent can't feel

A sandbox checks the code. It can't tell you what the change does under load.

What the agent gets today

Consequences, after the fact

Blind to consequences

Cost

What does the fleet cost overnight?

Latency

Does the queue drain under load?

Failure

What happens when a node disappears?

The missing layer

Between the plan and the merge

Validation checks the config. Modeling shows what it does.

```
# before merge
- behavior: unknown
+ behavior: modeled
```

Reads the pull request, models the resources it would create, and runs the workload against that model.

No environment. No shared staging cluster.

01 **Plan**

The agent writes the change.

↓

02 **Model**

Simulate what the change does.

↓

03 **Merge**

Ship knowing the cost, latency, and failure behavior.

What the run reports

1 MODELED, 4 ESTIMATED

Under 300 RPS of load, and with one node taken away.

Dimension	What you fed in	Value	Label
Cost	n1-standard-1 → e2-standard-2 × 2 nodes	**+$50.22/month ($0.73 → $0.80/hr)**	MODELED
Peak CPU	Same change at 300 RPS	**93.5% → 66.2%**	ESTIMATE
p99 latency	Same change at 300 RPS	**373 ms → 277 ms**	ESTIMATE
Error rate	Same change at 300 RPS	**6.00% → 0.00%**	ESTIMATE
One node lost	Same change, 1 of 2 nodes removed	**p99 337 → 264 ms, errors 10% → 6.75%**	ESTIMATE

Scope: request-serving compute and Kubernetes across the major clouds, GPU pools, and node cost from the actual shapes. Memory-behavior profiles: Cloud Run and GKE.

Cost is modeled; CPU, latency, error rate and node-loss figures are reference estimates.

The half that matters more

What it doesn't model yet

When the model can't answer, it tells you instead of guessing quietly.



Read the label before you quote a number

MODELED

The cost is modeled from list price. Your actual bill depends on commitments and discounts the model doesn't see.

Be wrong before the change is live, not after.



Not a guarantee

A planning estimate, not a guarantee. A real environment has the final word.

Thank you

If your team ships infrastructure as code every week, send one real pull request from your own repo.

One PR.
Blunt feedback.
No contract.

Kevin Brown · Cloud World Model
kevin@canvascloud.ai

cloudworldmodel.ai/teams

Setup guide: www.cloudworldmodel.ai/docs/github-pr-reports