

# Building Robust GPU Performance Validation into CI/CD Pipelines for AI Infrastructure

Presented By : Mohit Gupta

Intel Inc

Conf42 DevOps 2026.



## INTRODUCTION

# The Challenge: GPU Performance at Scale

Modern AI infrastructure demands systematic performance validation integrated directly into deployment workflows. Traditional post-deployment testing creates expensive feedback loops that slow innovation and increase costs.

Contemporary DevOps practices require shift-left approaches—detecting GPU performance bottlenecks during continuous integration before production deployment, not after.



# Why Performance Validation Matters

## Cost Impact

GPU resources represent significant infrastructure investment. Performance regressions directly translate to wasted compute cycles and increased operational expenses.

## Reliability Requirements

AI/ML workloads require predictable performance characteristics. Unexpected degradation impacts model training times, inference latency, and service-level agreements.

## Early Detection

Catching performance issues in CI/CD prevents production incidents, reduces mean time to resolution, and maintains deployment velocity for platform teams.

## FRAMEWORK OVERVIEW

# Comprehensive GPU Validation Framework

Our framework embeds GPU microarchitecture validation into automated CI/CD workflows, specifically addressing compute-bound versus memory-bound workload characteristics critical for infrastructure reliability.

The approach establishes quantitative classification criteria, automated performance gates, and observable metrics that enable platform engineering teams to maintain high-performance standards throughout the deployment lifecycle.

# Understanding Workload Characteristics



## Compute-Bound Workloads

Characterized by high arithmetic logic unit (ALU) utilization with minimal memory system pressure:

- ALU utilization exceeding 70%
- Memory stall cycles below 20%
- Dominated by mathematical operations



## Memory-Bound Workloads

Limited by data transfer capabilities rather than computational throughput:

- Cache miss rates above 15%
- Memory bandwidth utilization exceeding 60%
- Sensitive to access patterns and coalescing

# Classification Criteria: Quantitative Metrics

## Compute-Bound Indicators

- **ALU Utilization:** Greater than 70% sustained usage
- **Memory Stall Cycles:** Less than 20% of execution time
- **Cache Hit Rates:** Typically 85% or higher for L1

## Memory-Bound Indicators

- **Cache Miss Rate:** Above 15% for working sets
- **Bandwidth Utilization:** Exceeding 60% of theoretical peak
- **Coalescing Efficiency:** Critical for structured vs. irregular access



## VALIDATION METHODOLOGY

# Systematic Benchmark Validation

01

## Dense Linear Algebra Operations

Matrix multiplication, factorization, and vector operations representing compute-intensive AI workloads.

02

## Irregular Memory Access Patterns

Graph traversal, sparse operations, and random access patterns common in recommendation systems and knowledge graphs.

03

## Hybrid Execution Scenarios

Mixed workload validation combining compute and memory operations to identify interaction effects.

# Performance Analysis: Compute-Bound Results

85%

Sustained ALU Utilization

Achieved across representative dense linear algebra benchmarks

85%

L1 Cache Hit Rate

Demonstrates effective data locality in compute-intensive operations

<20%

Memory Stall Cycles

Minimal memory system bottlenecks during computation phases

# Performance Analysis: Memory-Bound Results

## Bandwidth Utilization

Memory-bound workloads reached **75% bandwidth utilization**, demonstrating significant memory system pressure characteristic of data-intensive AI operations.

## Coalescing Efficiency Variance

- **Structured Access:** 90% coalescing efficiency for sequential patterns
- **Irregular Patterns:** 45% efficiency revealing optimization opportunities



# Critical Discovery: Mixed Workload Interactions

- ❑ **Key Finding:** Concurrent execution of compute-bound and memory-bound workloads reveals performance degradation that isolated testing consistently misses.

1

## Isolated Testing

Individual workload characterization provides baseline performance expectations

2

## Interaction Effects

Resource contention and cache pressure during concurrent execution changes performance profiles

3

## Realistic Validation

Production environments require mixed workload testing to capture true performance characteristics



## IMPLEMENTATION

# Implementing Automated Validation as Code



## Define Performance Gates

Establish quantitative thresholds for ALU utilization, memory bandwidth, cache metrics, and coalescing efficiency within CI/CD configuration.



## Automate Benchmark Execution

Integrate representative workload suites into continuous integration workflows with consistent execution environments.



## Capture Observable Metrics

Collect GPU performance counters, memory access patterns, and utilization data for SRE observability and trend analysis.

# Platform Engineering Capabilities



## Prevent Performance Regressions

Automated gates block deployments that fail performance criteria, maintaining consistent service quality and preventing production incidents.



## Continuous Performance Visibility

Trend analysis and historical comparison enable proactive identification of performance drift and capacity planning.



## Shift-Left Optimization

Early detection reduces feedback loop time, accelerates development velocity, and minimizes cost of performance fixes.

# Building SRE Observability

## Key Observable Metrics

- GPU utilization trends across deployments
- Memory bandwidth saturation indicators
- Cache performance degradation alerts
- Workload classification distributions

These metrics enable SRE teams to establish service-level objectives, monitor infrastructure health, and respond proactively to performance anomalies.

# Key Takeaways for Platform Teams

## 1 Embed Validation in CI/CD

Shift-left performance testing prevents costly production issues and accelerates deployment confidence for AI infrastructure.

## 2 Classify Workload Characteristics

Quantitative metrics distinguish compute-bound from memory-bound workloads, enabling targeted optimization strategies.

## 3 Test Mixed Workload Scenarios

Interaction effects during concurrent execution reveal performance characteristics that isolated testing cannot capture.

## 4 Build Observable Systems

Performance metrics as code creates foundation for SRE practices, trend analysis, and continuous improvement.

# Thank You!

Mohit Gupta | Intel Inc

Conf42 DevOps 2026.