

# Building Real-Time AI Fraud Prevention Systems

for Gig Platforms in Rust

Prabhakar Singh, Senior Software Engineer at Meta



“Secure  
Your  
Hustle”

# Agenda

## The Gig Economy Fraud Challenge

Unique challenges, millisecond decision requirements, and why traditional solutions fall short

## Why Rust for Fraud Prevention

Memory safety, performance, and fearless concurrency advantages for security-critical systems

## Multi-Layered Architecture

Building with Tokio, Candle ML, Petgraph and implementing stream processing, zero-copy deserialization

## Real-World Results & Future Applications

Case studies from food delivery and ride-sharing platforms, plus emerging applications

# The Gig Economy Fraud Challenge

## Millisecond Decision Windows

Fraud prevention systems must approve/reject transactions within 50-100ms to maintain seamless user experience

## Decentralized Operations

Widely distributed workforce and customers create complex geographic fraud patterns

## Instant Payment Cycles

Immediate cashouts mean potential fraud losses cannot be recovered post-transaction

## AI Processing Requirements

ML models require high-throughput processing of thousands of transactions per minute

Legacy systems built with Python, Java, or Node.js often struggle with these requirements, creating painful tradeoffs between speed and accuracy.



# Common Fraud Patterns in Gig Platforms



- **Account Takeover:** Legitimate accounts compromised through credential stuffing
- **Synthetic Identity:** Fabricated identities created by combining real and fake information
- **Worker-Customer Collusion:** Coordinated fraud between delivery drivers and customers
- **GPS Spoofing:** Falsifying location data to manipulate ride/delivery assignments
- **Payment Method Fraud:** Using stolen payment methods for immediate cashouts

These patterns evolve rapidly, requiring systems that can adapt quickly to new fraud vectors while maintaining high performance.

# Why Rust for Fraud Prevention?



## Memory Safety Without Garbage Collection

Eliminates entire classes of bugs (buffer overflows, use-after-free) that can compromise fraud systems without runtime penalty



## C++-Level Performance


Zero-cost abstractions and LLVM-powered optimizations deliver processing speeds rivaling C++ with development velocity closer to Python



## Fearless Concurrency

Ownership and borrowing model prevents data races in concurrent processing, critical for high-throughput fraud detection

# Real-World Results: Performance Comparison



## Key Metrics from Production

- Rust achieves 93% of C++'s raw performance
- 6.2x faster than Java implementation
- 12.4x faster than Python with TensorFlow
- Consistent p99 latency under 35ms
- 90% reduction in CPU utilization compared to Python

Data from benchmark tests on a major food delivery platform processing 50,000+ orders per hour.

# Multi-Layered Rust Architecture for Fraud Prevention

## 1 Real-Time Stream Processing with Tokio

Asynchronous processing of transaction streams using Tokio and Futures, drastically reducing detection latency

## 2 Rule-Based Detection Layer

High-speed pattern matching using Rust's match expressions and custom DSLs

## 3 Statistical Anomaly Detection

Time-series analysis with statistical models implemented in pure Rust

## 4 ML Classification with Candle

Fast ML inference using Candle for deep learning models with CUDA acceleration

## 5 Graph Analysis with Petgraph

Network analysis for detecting coordinated fraud rings

# Rust-Powered Innovations



## Async Stream Processing

Tokio and Futures enable processing thousands of transactions per second with minimal latency

```
// Transaction processing with Tokio
let mut stream = StreamExt::throttle(
    transactions, Duration::from_millis(1)
);

while let Some(tx) = stream.next().await
{
    let verdict = detect_fraud(tx).await?;
    decisions.push(verdict).await?;
}
```



## Zero-Copy Deserialization

Serde implementations with zero-copy parsing for high-throughput JSON/Protobuf data

```
// Zero-copy deserialization with Serde
#[derive(Deserialize)]
struct Transaction<'a> {
    #[serde(borrow)]
    user_id: &'a str,
    amount: f64,
    #[serde(borrow)]
    device_id: &'a str,
}
```



## WASM-Compiled Modules

Secure, portable fraud detection logic that runs consistently across diverse environments



# ML Inference with Candle

## Why Candle for ML in Fraud Detection?

- Pure Rust ML framework - memory safety throughout the stack
- CUDA and Metal acceleration for GPU inference
- Supports ONNX model imports from PyTorch/TensorFlow
- No Python dependencies in production
- Tight integration with Rust's type system



# Fraud patterns

## Graph Analysis for Fraud Networks

- Detecting Coordinated Fraud with Petgraph

Rust's Petgraph crate enables efficient graph algorithms for finding fraud networks:

- Connected component analysis reveals collusion between workers and customers
- Centrality measures identify key nodes in fraud networks
- Graph embeddings detect structural similarities in transaction patterns
- Temporal graph analysis tracks evolving fraud patterns over time

- Real-World Example: Driver-Customer Collusion

A food delivery platform detected a sophisticated fraud ring involving 32 delivery drivers and 18 customers placing fake orders. Graph analysis revealed the pattern when standard ML models missed it.

```
// Find suspicious components in the graph
let components = petgraph::algo::kosaraju_scc(&graph);
for component in components {
    if component.len() > 5 {
        analyze_potential_fraud_ring(component);
    }
}
```


# Privacy-Preserving Fraud Detection

## Homomorphic Encryption

Rust implementations of FHE libraries allow computation on encrypted data without decryption

## Secure Enclaves

Rust code running in TEEs provides hardware-level isolation for sensitive fraud logic



## Federated Learning

Distributed ML training across devices/regions without centralizing sensitive data

## Differential Privacy

Adding calibrated noise to protect individual privacy while preserving statistical utility

Rust's strong type system and zero-cost abstractions make it possible to implement these advanced privacy techniques without prohibitive performance penalties.

# Case Study: Major Ride-Sharing Platform



## Challenge

Platform was losing \$3.2M monthly to sophisticated fraud schemes including GPS spoofing and synthetic identity creation.

## Rust Solution

Replaced Python-based system with Rust microservices using:

- Tokio for async HTTP API and stream processing
- Custom Rust implementation of device fingerprinting
- WASM modules for distributed fraud rules
- Candle for ML model inference

## Results

87% reduction in fraud losses, 65% decrease in false positives, and 99.98% uptime over 12 months.



# Emerging Applications in Rust Fraud Prevention





# Key Takeaways

## Performance Matters

Rust delivers near-C++ performance with development velocity approaching higher-level languages, critical for millisecond-scale fraud decisions

## Safety Is Non-Negotiable

Memory safety guarantees and concurrency without data races eliminate entire classes of vulnerabilities in security-critical systems

## Ecosystem Is Maturing

Crates like Tokio, Candle, and Petgraph now provide production-ready foundations for sophisticated fraud prevention systems