

# Scaling Microservices

## Development A

### Shared-Cluster

## Deployment Platform

## for Faster, Cost-Efficient

## SaaS Delivery

By : Ravi Babu Dasari, NetApp Inc.,

Conf42 DevOps 2026

# The Microservices Promise and Challenge

## The Promise

Microservices architectures deliver unparalleled scalability and modularity, enabling rapid innovation and independent service deployment across modern SaaS platforms.

## The Challenge

However, they introduce complexity in environment management, testing workflows, and cross-team coordination that traditional infrastructure approaches struggle to address efficiently.



# Traditional Approach: The Bottleneck

## Dedicated Clusters

Provisioning separate Kubernetes clusters for every feature branch creates isolation but at tremendous cost.


## Resource Waste

Excessive infrastructure usage with underutilised compute resources scattered across numerous environments.

## Slow Provisioning

Long wait times for cluster creation become critical bottlenecks that slow delivery pipelines significantly.

# Traditional Architecture



# Why Canary Deployments Fall Short for Feature Branch Testing

- **Designed for Production Rollouts, Not Development**
- **Mostly limited to just one canary deployment at a time**
- **Complications with Stateful Services**
- **Limited Isolation in Shared Clusters**
- **Better Alternative: Namespace-Based Smart Isolation for any number of feature branches**

# Introducing SDDP

The Scalable Development Deployment Platform



## Namespace

### Isolation

Leverage Kubernetes namespace boundaries to provide secure, isolated environments within shared clusters.

## Dynamic Routing

Implement intelligent ingress routing that directs traffic to feature-specific microservices seamlessly.

## Shared

### Infrastructure

Enable feature services to run alongside production-grade components without full-stack duplication.

# Deploying a feature branch to a namespace with dynamic routing

YAML

X




Copy

```
# values-feature-mybranch.yaml (Helm override for feature branch)
replicaCount: 1
image:
  tag: "mybranch-abc123"

# Auto-generated by SDDP portal
namespace: feature-mybranch

ingress:
  enabled: true
  host: mybranch.myapp.dev.example.com
  annotations:
    kubernetes.io/ingress.class: nginx
    nginx.ingress.kubernetes.io/rewrite-target: /
```

# Proposed Architecture



# SDDP Architecture

## Overview

### Management Portal


1 Web-based interface for orchestrating deployments, managing environments, and providing visibility across teams.

### Service Layer

2 Integrates seamlessly with CI/CD pipelines through Helm-based automation and declarative configuration management.

### Infrastructure Layer

3 Kubernetes foundation optimised for multi-tenancy, efficient resource allocation, and horizontal scaling capabilities.



# Demonstrate namespace isolation with NetworkPolicy

YAML

X ⌂ Copy

```
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: deny-cross-namespace
  namespace: feature-mybranch
spec:
  podSelector: {}
  policyTypes:
    - Ingress
    - Egress
  ingress:
    - from:
        - namespaceSelector:
            matchLabels:
              kubernetes.io/metadata.name: feature-mybranch # Only same namespace
        - podSelector: {} # Allow from baseline services if needed
  egress:
    - to:
        - namespaceSelector:
            matchLabels:
              app: baseline-services
```

# Core Benefits of the Platform



## Streamlined Development

Accelerated workflows that maintain strict isolation whilst reducing operational overhead.

## Cost Efficiency

Dramatic reduction in infrastructure spending through intelligent resource sharing and utilisation.

## Environment Consistency

Unified configuration management ensures parity between development, staging, and production environments.

# Security Without Compromises

## Role-Based Access Control

Granular permissions ensure teams access only their designated namespaces and resources, maintaining strict security boundaries.

## Centralised Secret Management

Integrated secret management solutions provide encrypted storage and controlled access to sensitive configuration data.

## Network Policies

Kubernetes network policies enforce traffic segregation between namespaces, preventing unauthorised service-to-service communication.

# Handling Outliers: Challenges with Stateful Services

Core

## Challenge

Stateful services (e.g., databases, caches, blob storage) introduce side effects when shared across namespaces in Kubernetes.

## Risks

Data corruption, inconsistent states, or interference between feature branches/testing environments.

## Outliers in Microservices

Most services are stateless and easy to isolate, but stateful ones and complex dependencies require special handling.

**Goal:** Balance isolation for safe testing with resource efficiency in shared clusters.

# Strategies for Isolation and Replication

## Tools & Techniques

Kubernetes namespaces for resource segregation.  
Ingress/service meshes for traffic routing (e.g., connection strings).  
Cache isolation: Namespace-scoped keys or dedicated instances.

## Automation with AI

Analyze code changes, recommend services to replicate, and auto-generate Kubernetes manifests (YAML for namespaces/ingress).


## Change Analysis

Scan feature branch diffs to detect impacts on stateful components.

Examples:

- New endpoint (no schema change) → Share existing database.
- Schema migration → Replicate affected database in namespace.

# Analyze with the help of AI



# AI-driven change analysis logic (pseudo-code in Python)

```
Python
def analyze_changes(diff):
    stateful_impacts = []

    if any('migration' in file or '.sql' in file for file in diff):
        stateful_impacts.append({
            "service": "postgresql",
            "action": "replicate",
            "reason": "schema migration detected"
        })

    if any('redis' in file and 'KEYS' in content for file in diff):
        stateful_impacts.append({
            "service": "redis",
            "action": "namespace-scoped-keys",
            "reason": "direct key access - risk of collision"
        })

    return stateful_impacts

# SDDP uses this to decide what to replicate vs. share
```

# Dynamic connection string routing via environment variable injection

```
YAML  
  
# Deployment snippet auto-injected by SDDP  
env:  
  - name: DATABASE_URL  
    value: "postgres://baseline-db.prod.svc.cluster.local:5432/app"  
  - name: REDIS_URL  
    value: "redis://redis-feature-mybranch:6379" # Replicated instance  
  - name: CACHE_PREFIX  
    value: "mybranch:" # Namespace-scoped keys when sharing Redis
```

# ResourceQuota to prevent noisy neighbor issues

```
YAML  
  
apiVersion: v1  
kind: ResourceQuota  
metadata:  
  name: feature-quota  
  namespace: feature-mybranch  
spec:  
  hard:  
    requests.cpu: "4"  
    requests.memory: 8Gi  
    limits.cpu: "8"  
    limits.memory: 16Gi  
    pods: "20"
```

# Integration with CI/CD Workflows

1

## Commit Trigger

Developer pushes code to feature branch, initiating automated pipeline.

2

## Build & Package

CI system builds container images and packages Helm charts automatically.

3

## Deploy to SDDP

Platform provisions namespace and deploys services with dynamic routing.

4

## Test & Validate

Automated tests run against isolated environment before promotion.

# Measured Impact and Outcomes

- **Faster Environment Creation**

Reduction in time from request to usable development environment

- **Shorter Testing**

~~Cycles~~ Decrease in end-to-end testing duration through parallel execution

- **Resource Utilisation**

Improvement in compute resource efficiency across infrastructure

Engineering teams adopting SDDP observed substantial improvements across all key delivery metrics whilst significantly reducing operational costs.

# Enhanced Team Collaboration

## Development Teams

- Self-service environment provisioning
- Rapid iteration without infrastructure delays
- Clear visibility into deployment status
- Consistent tooling across all environments

## DevOps Teams

- Centralised management and monitoring
- Reduced operational burden and toil
- Standardised deployment patterns
- Improved resource governance



# Key Technical Considerations

## Resource Quotas

Implement namespace-level resource limits to prevent resource exhaustion and ensure fair allocation across teams.

## Ingress Strategy

Design dynamic routing patterns that support multiple ingress controllers and custom domain configurations.

## Data Management

Establish clear patterns for database provisioning, test data management, and state isolation between environments.

## Observability

Deploy comprehensive logging, metrics, and tracing infrastructure to maintain visibility across shared resources.

# Your Path to Adoption

## Phase 1: Foundation

Establish shared cluster infrastructure and core namespace isolation patterns.

## Phase 3: Security

Implement RBAC, network policies, and secret management solutions.



## Phase 2:

### Automation

Build CI/CD integration and self-service deployment capabilities.

## Phase 4: Scale

Optimise resource allocation and expand to additional teams and services.

# Transform Your Development Workflow

- **Accelerate Delivery**

Reduce environment provisioning from hours to minutes

- **Scale Confidently**

Support growing teams without proportional infrastructure growth

- **Optimise Costs**

Dramatically decrease infrastructure spending through sharing



# Real-World Success: Six Years of Proven Results

- **The Bottleneck**

Manual cluster setup for each feature branch wasted hours daily, severely hindering development progress.

- **The Solution**

A shared cluster with a dedicated portal streamlined environment management, delivering six years of continuous success.

- **Key Outcomes**

Deployment times dropped from hours to minutes, accelerating development, boosting productivity, and improving team morale and system reliability.




# The Deployment Portal: Your Command Center

- Automated Sandbox Creation
  - Instantly create sandboxes for any feature branch, automating deployment and baseline selection for rapid testing.
- Effortless Collaboration
  - Share sandbox links for quick team reviews and demos, accelerating feedback and alignment.
- AI-Powered Intelligence
  - AI analyzes code changes, predicts side effects, and suggests optimal configurations for efficient, intelligent deployments.



# Transformative Benefits and ROI

- Cost Reduction Infrastructure cost savings.
- Setup Time & Deployment time reduction.
- Engineers Supported Engineers enabled at scale.
- Operational Load Consistent operational overhead.



# Start Your Transformation Today

- Revolutionize Development

Redefine microservices development using robust infrastructure and AI-driven tools for enhanced efficiency and innovation.

- Proven Impact

This methodology delivers substantial cost reductions, faster delivery cycles, and enhanced collaboration.

- Start Your Journey

Initiate transformation with a strategic pilot program, establish a baseline, and scale success across your organization.

Join our community and share your transformation journey to shape the future of development.

# Demo

Source code: <https://github.com/RaviDasari/sddp>

# Thank You! Questions?

---

**Ravi Babu Dasari**  
**NetApp Inc.,**  
**IncConf42 DevOps 2026.**



[linkedin.com/in/ravi-babu-dasari-6528621b](https://www.linkedin.com/in/ravi-babu-dasari-6528621b)



[@ravibabudasari](https://twitter.com/ravibabudasari)

[github.com/RaviDasari](https://github.com/RaviDasari)