


Operationalizing LLMs at Scale Reliability, Resilience, and Cognitive Commerce in Grocery Retail

By : Sanjay Basu

- TCS
- Conf42 Site Reliability Engineering 2026



The Shift: From Experiment to Mission-Critical



LLMs are no longer proof-of-concepts tucked away in innovation labs. In grocery retail and e-commerce, they are becoming **load-bearing infrastructure** processing substitutions, personalising recommendations, and orchestrating fulfilment in real time.

For SREs, this means probabilistic, multimodal AI systems must now satisfy the same expectations as any Tier-1 service: **availability, latency, cost efficiency, and compliance.**



Why LLMs Break Conventional SRE Assumptions

Non-Deterministic Outputs

The same prompt can yield different results across calls invalidating traditional correctness checks.

Token-Level Latency Variability

Response time scales with output length, not just input complexity making P99 SLOs extremely hard to enforce.

Evolving Context Dependencies


Prompts, embeddings, and context windows change continuously, introducing silent regressions that bypass standard CI/CD gates.

The Retail Stakes Are Uniquely High



Peak Traffic Spikes

Flash sales, promotions



Product Substitutions

Out-of-stock AI decisions



Personalisation at Scale

Real-time recommendations per user



Inventory Accuracy

Shelf vs system discrepancy detection

In grocery retail, **model failure is not an abstract engineering problem** it translates directly into wrong substitutions, lost baskets, and eroded customer trust.

Peak traffic events flash sales, holiday surges amplify every latency and accuracy weakness. Reliability engineering must account for these operational realities from day one, not as afterthoughts.

Redefining SLOs for Probabilistic Systems

Traditional SLOs measure binary correctness and latency. For LLMs, you need a richer contract.

Semantic Accuracy SLO

% of responses that pass automated semantic evaluation against ground-truth intent not just format validation.

Time-to-First-Token (TTFT)

Separate latency budget for stream initiation vs. full completion. Retail UX demands TTFT under 300ms.


Hallucination Rate Budget

Acceptable error rate for factually incorrect outputs monitored per model version and prompt template, not globally.

Cost-per-Inference SLO

Token spend as a first-class SLO. Runaway context windows can double infrastructure cost overnight.

Architectural Resilience Patterns



Each layer acts as a blast-radius limiter. A guardrail failure should never propagate to a customer-facing outage — the architecture must absorb it silently and degrade predictably, not catastrophically.

Hybrid Decision Layers: AI + Rules

Pure LLM decision-making is inappropriate for high-stakes retail actions such as **out-of-stock substitutions, pricing overrides, or fraud flags**. A hybrid architecture enforces deterministic guardrails around probabilistic reasoning.

Why This Matters

In regulated retail alcohol, tobacco, pharmacy an unconstrained model output can constitute a compliance breach. Hybrid layers are your legal and operational safety net.

→ Rule Engine First

Hard business constraints evaluated before any LLM call
category substitution rules,
allergen restrictions, price caps.

→ LLM as Reasoning Layer

Model operates within a bounded decision space defined by rules
reducing hallucination surface area significantly.

→ Post-Inference Validation

Output re-checked against rule engine before surfacing to customer or downstream systems.



∞ OBSERVABILITY

Observability Frameworks for AI Systems

- **Prompt Telemetry**

Log every prompt template version, variable injection, and token count. Treat prompts as versioned code artefacts in your observability pipeline.

- **Hallucination Detection**

Embedding-based similarity checks between model output and retrieval context. Flag low-confidence responses before they surface to users.

- **Semantic Evaluation Pipelines**


Automated LLM-as-judge scoring on sampled production traffic measuring coherence, relevance, and policy compliance continuously.

- **Drift Monitoring**

Track semantic distribution shifts across model versions, prompt updates, and embedding refreshes not just accuracy metrics.

Cost-Aware Scaling: Token Spend as Infrastructure

— Optimised (tokens/req) — Baseline (tokens/req)



Token spend is a **first-class SRE concern**. Unoptimised context windows can inflate inference costs 3–4× without any change in model version.

- Prompt compression and context pruning as platform primitives
- Tiered model routing: small models for simple intents, large models for complex reasoning
- Cached embedding lookups to avoid redundant retrieval on repeated queries

LLMs Across the Retail Value Chain



Customer Experience

Semantic search, conversational shopping assistants, and real-time substitution recommendations.



Merchandising

Demand signal interpretation, assortment optimisation, and automated promotional copy generation.



Supply Chain

Supplier communication automation, anomaly detection in fulfilment data, and inventory forecasting narratives.


Agentic LLMs in Logistics & Last-Mile Delivery

Agentic architectures where LLMs autonomously plan, invoke tools, and chain decisions introduce a **new category of operational risk** in fulfilment centres and last-mile routing.

Autonomous reasoning must be **bounded by strict reliability constraints**: action budgets, rollback triggers, and human escalation thresholds enforced at the orchestration layer not left to the model's discretion.



Incident Response for AI-Induced Failures



Security, Compliance & Guardrail Enforcement

Input Guardrails

Prompt Injection Detection

Pattern classifiers on all user-supplied input before LLM invocation.

PII Scrubbing

Named entity recognition strips personal data from context windows before inference.

Output Guardrails

Regulatory Filters

Hard blocks on alcohol, pharmacy, and age-restricted product recommendations without verified eligibility.

Toxicity & Bias Screening

Policy classifiers on every model output before it reaches any customer-facing surface.

A Pragmatic Framework for Cognitive Commerce

- **Redefine SLOs**

Extend your SLO framework to cover semantic accuracy, hallucination rate, TTFT, and cost-per-inference as first-class reliability signals.

- **Architect for Degradation**

Every LLM call must have a deterministic fallback. Resilience is designed in not bolted on after the first production incident.

- **Instrument Everything**

Prompt telemetry, semantic evaluation, drift monitoring, and cost tracking are your observability baseline not optional enhancements.

- **Bound Agentic Autonomy**

Action budgets, rollback triggers, and human escalation thresholds must be enforced at the orchestration layer for every agentic workflow.

Thank You!
Discussions?