FROM ALERTS TO ANSWERS: MODERN SRE WITH OBSERVABILITY, AI AND AGENTIC OPS

Presented by- Savi Grover

Conf42 Site Reliability Engineering (SRE) 2026

The Key Problem - Noise data rich + Information Poor

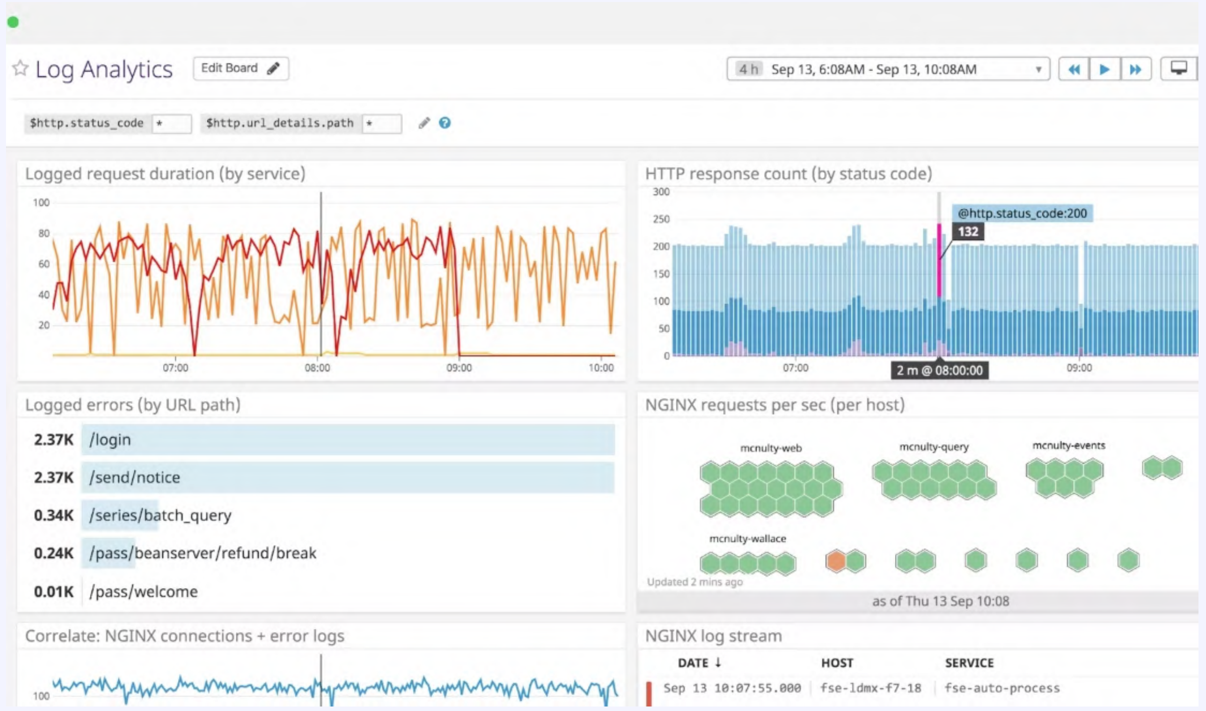

The "Data Rich, Information Poor" Paradox

- **The Shift:** We moved from monoliths to microservices, increasing telemetry volume by 100x.
- **The Fatigue:** SREs are drowning in **"High-Cardinality Noise"**—too many alerts, not enough context.
- **The Goal:** We don't need more alerts; we need **answers**.
- **The Bottom Line:** Traditional monitoring identifies the *symptom*, but SREs are still manually hunting for the *cause*.

Challenges with Current Observability

Why the "Old Way" is Failing

- **Siloed Data:** Metrics, logs, and traces exist in separate tools with no correlation.
- **The "Dashboard Mirage":** Looking at 50 dashboards during an outage increases cognitive load rather than reducing it.
- **Alert Fatigue:** Static thresholds lead to 80% "flapping" or non-actionable alerts.
- **Knowledge Leakage:** Incident resolution depends on "tribal knowledge" rather than encoded patterns.

Coping with Signal Bias & Noise

"Signal-to-Context" Framework

To handle huge signal bias, we shift focus from raw data to **meaningful events**:

- **Golden Signals Re-imagined:** Focus on Latency, Errors, Traffic, and Saturation, but through the lens of **Service Level Objectives (SLOs)**.
- **Distributed Tracing as the Backbone:** Using trace IDs to stitch together the user journey across the mesh.
- **Dimensionality Reduction:** Using AI to group 1,000 "Connection Refused" logs into 1 "Database Connectivity Cluster" event.

Handling Signal Bias & Noise

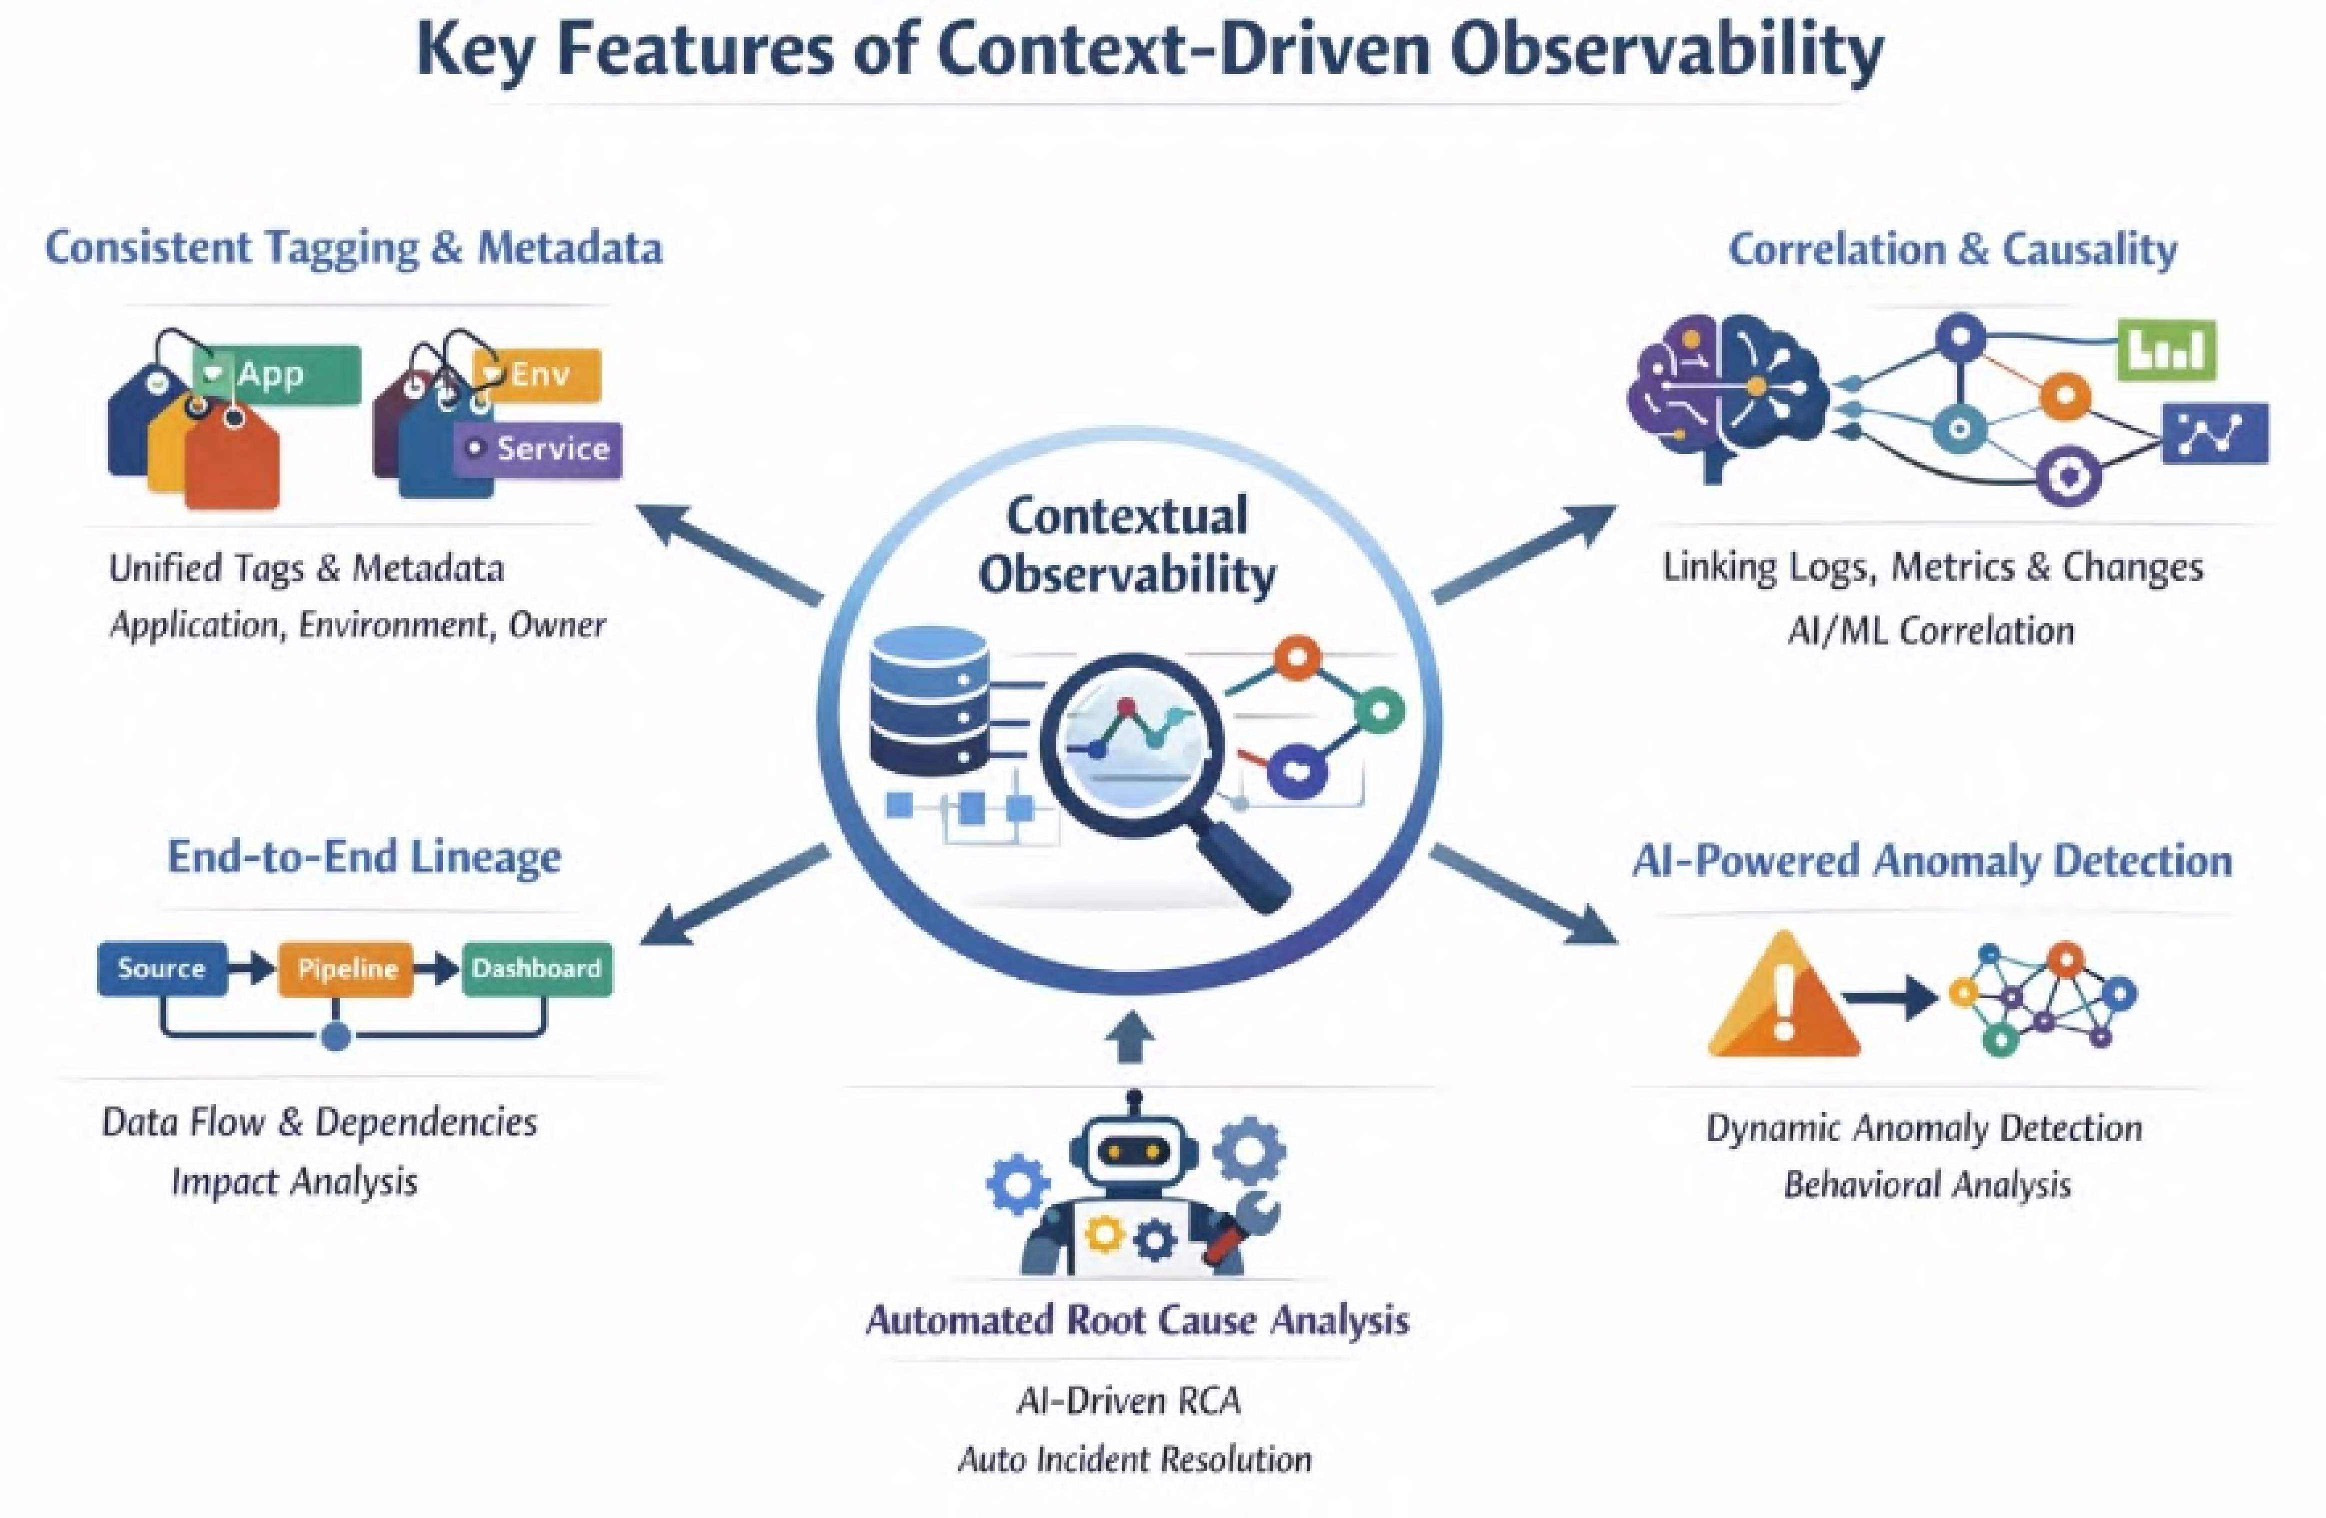

Key Features of Context-Driven Observability:

- **Consistent Tagging and Metadata:** Enriching telemetry data with uniform tags (e.g., application name, environment, service tier, owner).
- **Correlation and Causality:** Tools use AI/ML to automatically correlate unrelated data points into a coherent picture of a problem. This includes linking logs, metrics, and traces to specific code changes, infrastructure components, or business impacts.
- **End-to-End Lineage:** showing how data flows and which downstream assets are affected by an upstream issue. This helps with impact analysis and root cause identification.
- **AI-Powered Anomaly Detection:** Instead of relying solely on static thresholds, context-aware tools learn normal behavioral patterns and flag anomalies.
- **Automated Root Cause Analysis (RCA):** The ultimate goal of context is to expedite RCA. To go beyond detection, automatically triaging incidents and suggesting potential fixes.

Agentic Ops

What are AI Agents in SRE?

Unlike a script, an Agent can reason.

1. **Observe**: Agent detects a spike in 5xx errors.
2. **Analyze**: Agent queries traces and sees a specific downstream API is timing out.
3. **Act**: Agent checks the recent CI/CD logs, identifies a bad canary, and triggers a roll-back.
4. **Report**: Agent posts a summary in Slack: *"I rolled back v1.2 due to API latency; here is the trace evidence."*

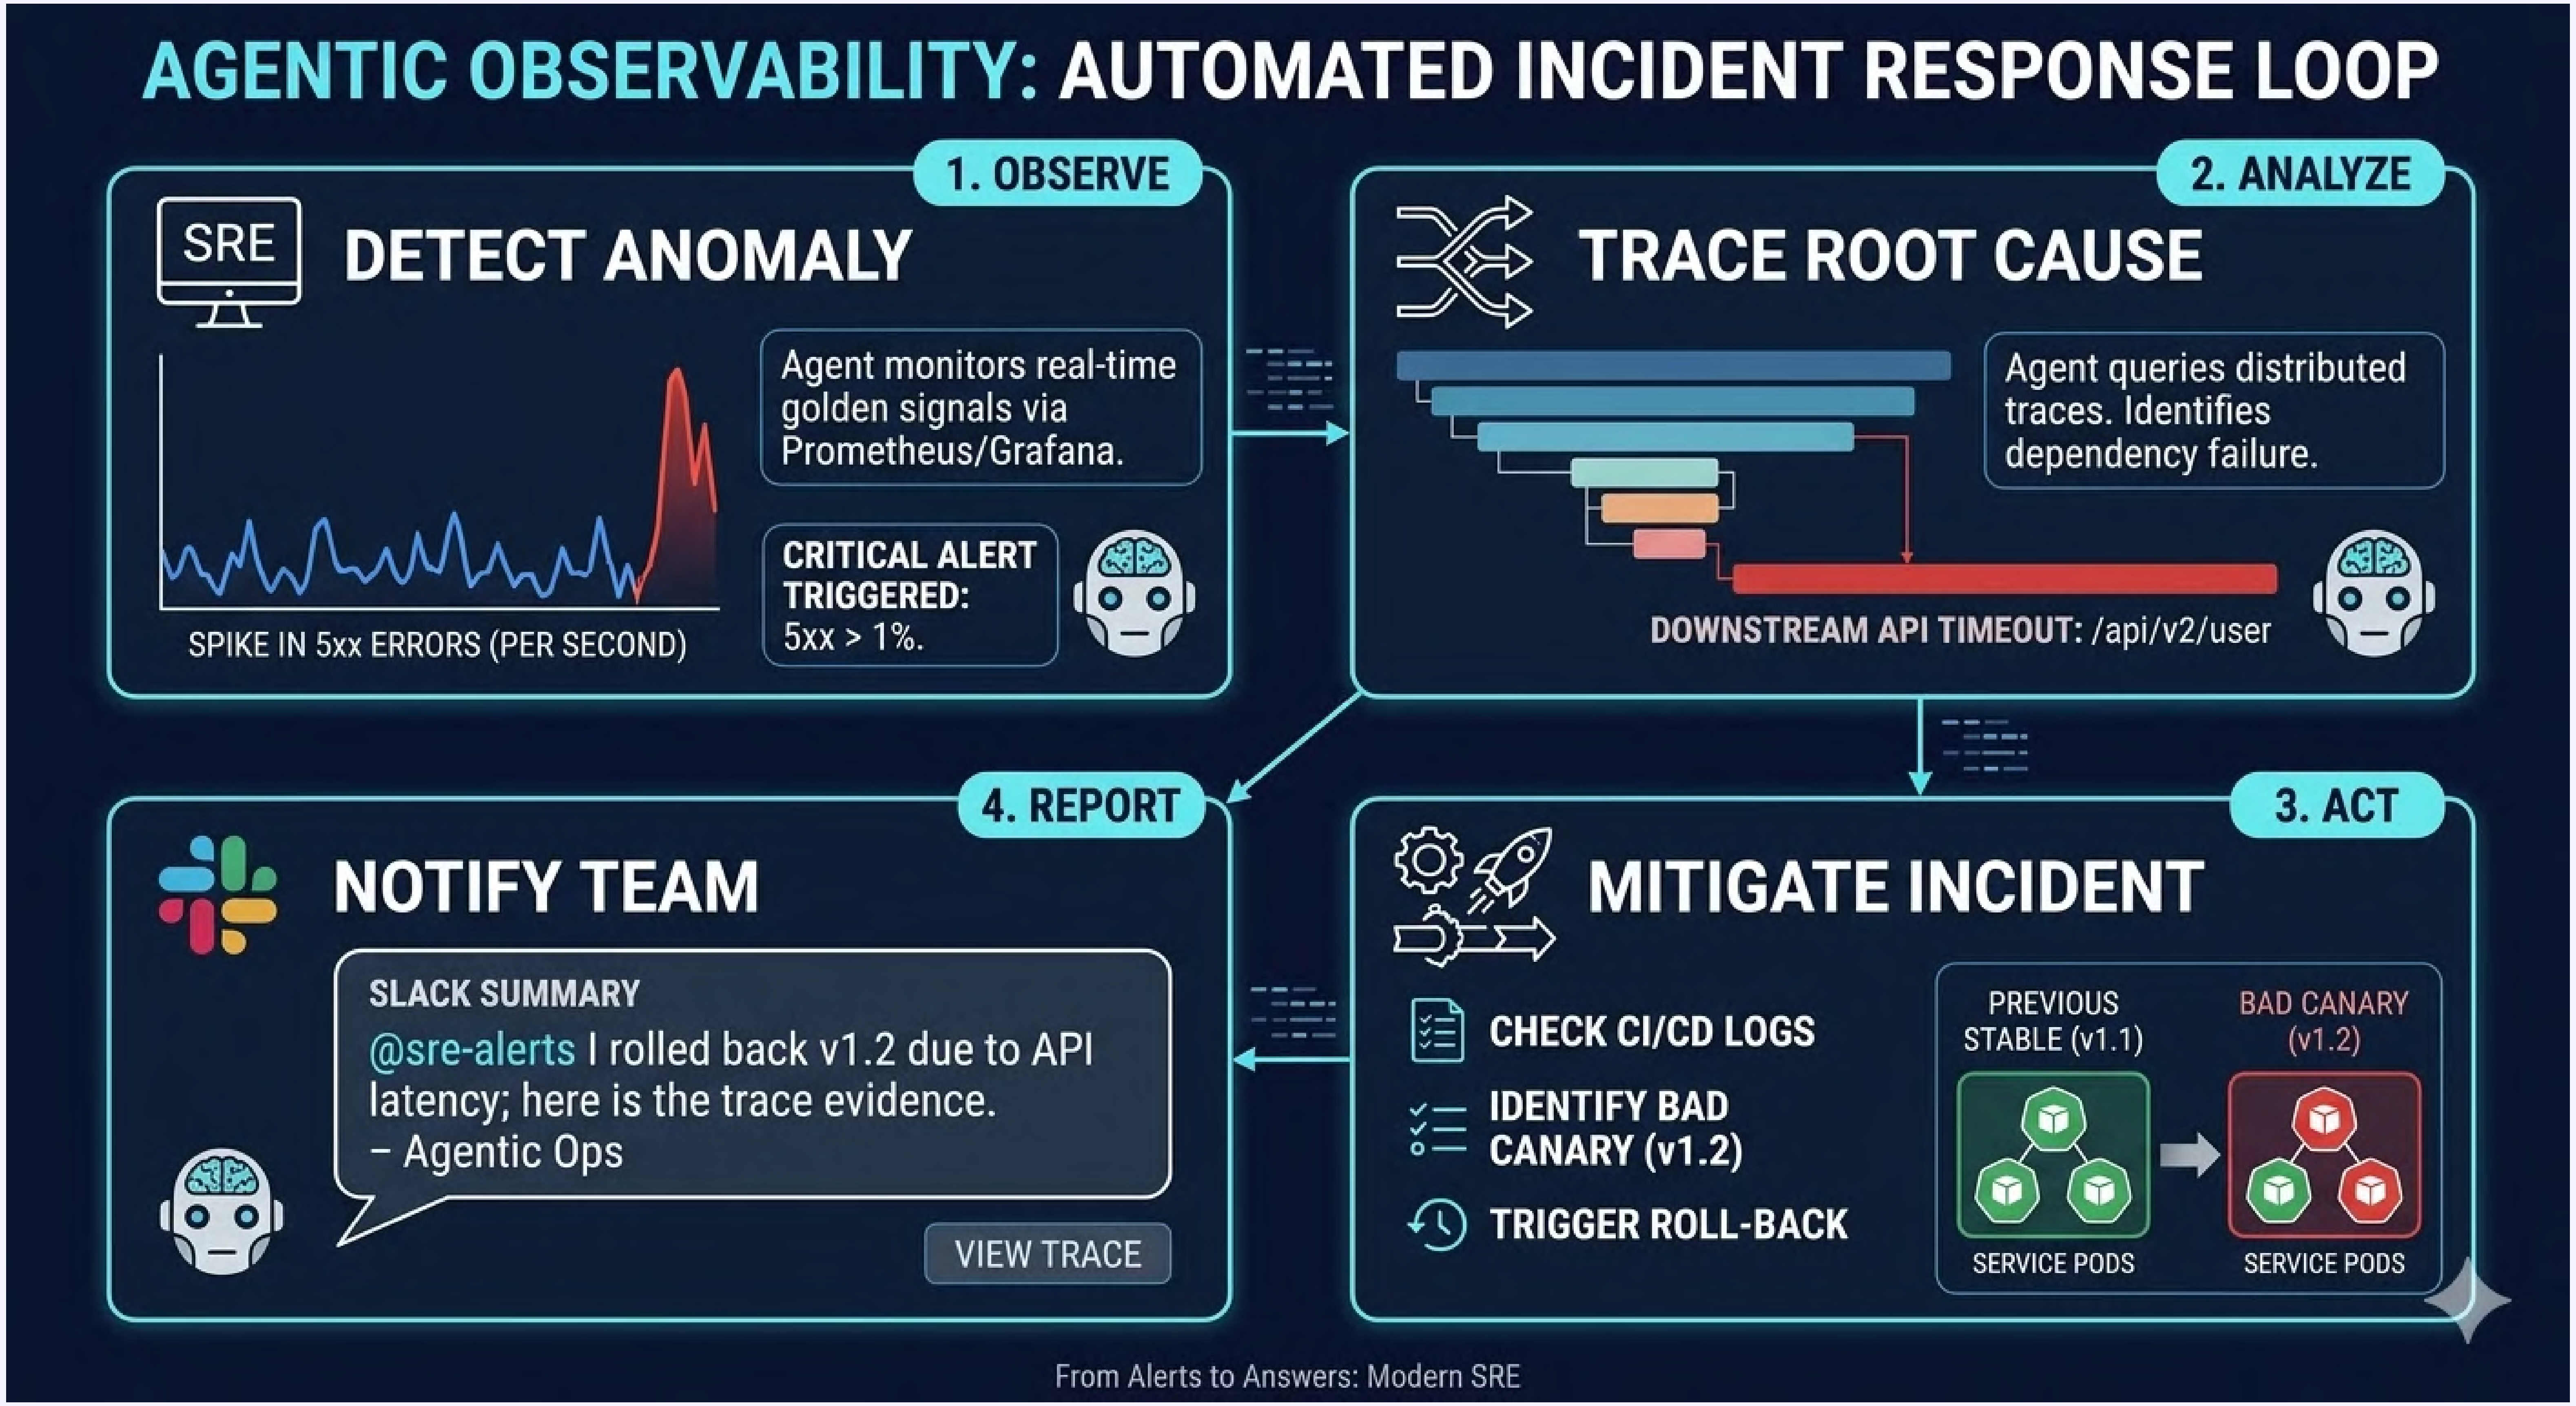
AGENTIC OBSERVABILITY: AUTOMATED INCIDENT RESPONSE LOOP
1. OBSERVE
SRE
DETECT ANOMALY
Agent monitors real-time golden signals via Prometheus/Grafana.
CRITICAL ALERT TRIGGERED: 5xx > 1%.
SPIKE IN 5xx ERRORS (PER SECOND)
2. ANALYZE
TRACE ROOT CAUSE
Agent queries distributed traces. Identifies dependency failure.
DOWNSTREAM API TIMEOUT: /api/v2/user
3. ACT
MITIGATE INCIDENT
CHECK CI/CD LOGS
IDENTIFY BAD CANARY (v1.2)
TRIGGER ROLL-BACK
PREVIOUS STABLE (v1.1)
BAD CANARY (v1.2)
SERVICE PODS
SERVICE PODS
4. REPORT
NOTIFY TEAM
SLACK SUMMARY
@sre-alerts I rolled back v1.2 due to API latency; here is the trace evidence.
– Agentic Ops
VIEW TRACE
From Alerts to Answers: Modern SRE

Thankyou

Pitch

Want to make a presentation like this one?

Start with a fully customizable template, create a beautiful deck in minutes, then easily share it with anyone.

Create a presentation (It’s free)